

Artificial Neural Networks




Data Base and Data Mining Group of Politecnico di Torino

Elena Baralis
Politecnico di Torino




Artificial Neural Networks

- Inspired to the structure of the human brain
 - Neurons as elaboration units
 - Synapses as connection network



Biological Neuron







Artificial Neural Networks

■ Different tasks, different architectures


image understanding: convolutional NN (CNN)




time series analysis: recurrent NN (RNN)



numerical vectors classification: feed forward NN (FFNN)




denoising: auto-encoders






Feed Forward Neural Network





Structure of a neuron

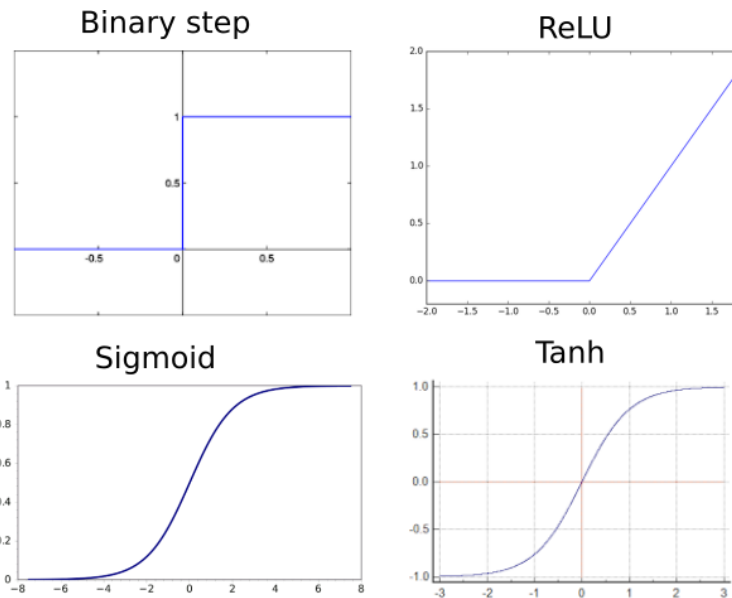




Activation Functions

■ Activation

- simulates biological activation to input stymula
- provides non-linearity to the computation
- may help to saturate neuron outputs in fixed ranges







Activation Functions

- Sigmoid, tanh
 - saturate input value in a fixed range
 - non linear for all the input scale
 - typically used by FFNNs for both hidden and output layers
 - E.g. *sigmoid* in output layers allows generating values between 0 and 1 (useful when output must be interpreted as likelihood)

Sigmoid



Tanh






Activation Functions

■ Binary Step

- outputs 1 when input is non-zero
- useful for binary outputs
- **issues:** not appropriate for gradient descent
 - derivative not defined in $x=0$
 - derivative equal to 0 in every other position


Binary step



■ ReLU (Rectified Linear Unit)

- used in deep networks (e.g. CNNs)
 - avoids vanishing gradient
 - does not saturate
- neurons activate linearly only for positive input

ReLU



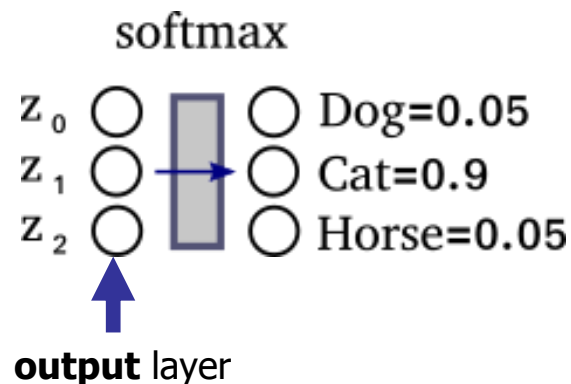


Activation Functions

■ Softmax

- differently to other activation functions
 - it is applied only to the **output** layer
 - works by considering **all the neurons** in the layer
- after softmax, the output vector can be interpreted as a discrete **distribution of probabilities**
 - e.g. the probabilities for the input pattern of belonging to each class

$$\text{softmax}(z_j) = \frac{e^{z_j}}{\sum_{i=0}^{N-1} e^{z_i}}$$





Building a FFNN

- For each node, definition of
 - set of weights
 - offset valueproviding the highest accuracy on the training data
- Iterative approach on training data instances



Building a FFNN

■ Base algorithm

- Initially assign random values to weights and offsets
- Process instances in the training set one at a time
 - For each neuron, compute the result when applying weights, offset and activation function for the instance
 - Forward propagation until the output is computed
 - Compare the computed output with the expected output, and evaluate error
 - Backpropagation of the error, by updating weights and offset for each neuron
- The process ends when
 - % of accuracy above a given threshold
 - % of parameter variation (error) below a given threshold
 - The maximum number of epochs is reached



Feed Forward Neural Networks

■ Strong points

- High accuracy
- Robust to noise and outliers
- Supports both discrete and continuous output
- Efficient during classification


■ Weak points

- Long training time
 - weakly scalable in training data size
 - complex configuration
- Not interpretable model
 - application domain knowledge cannot be exploited in the model



Convolutional Neural Networks

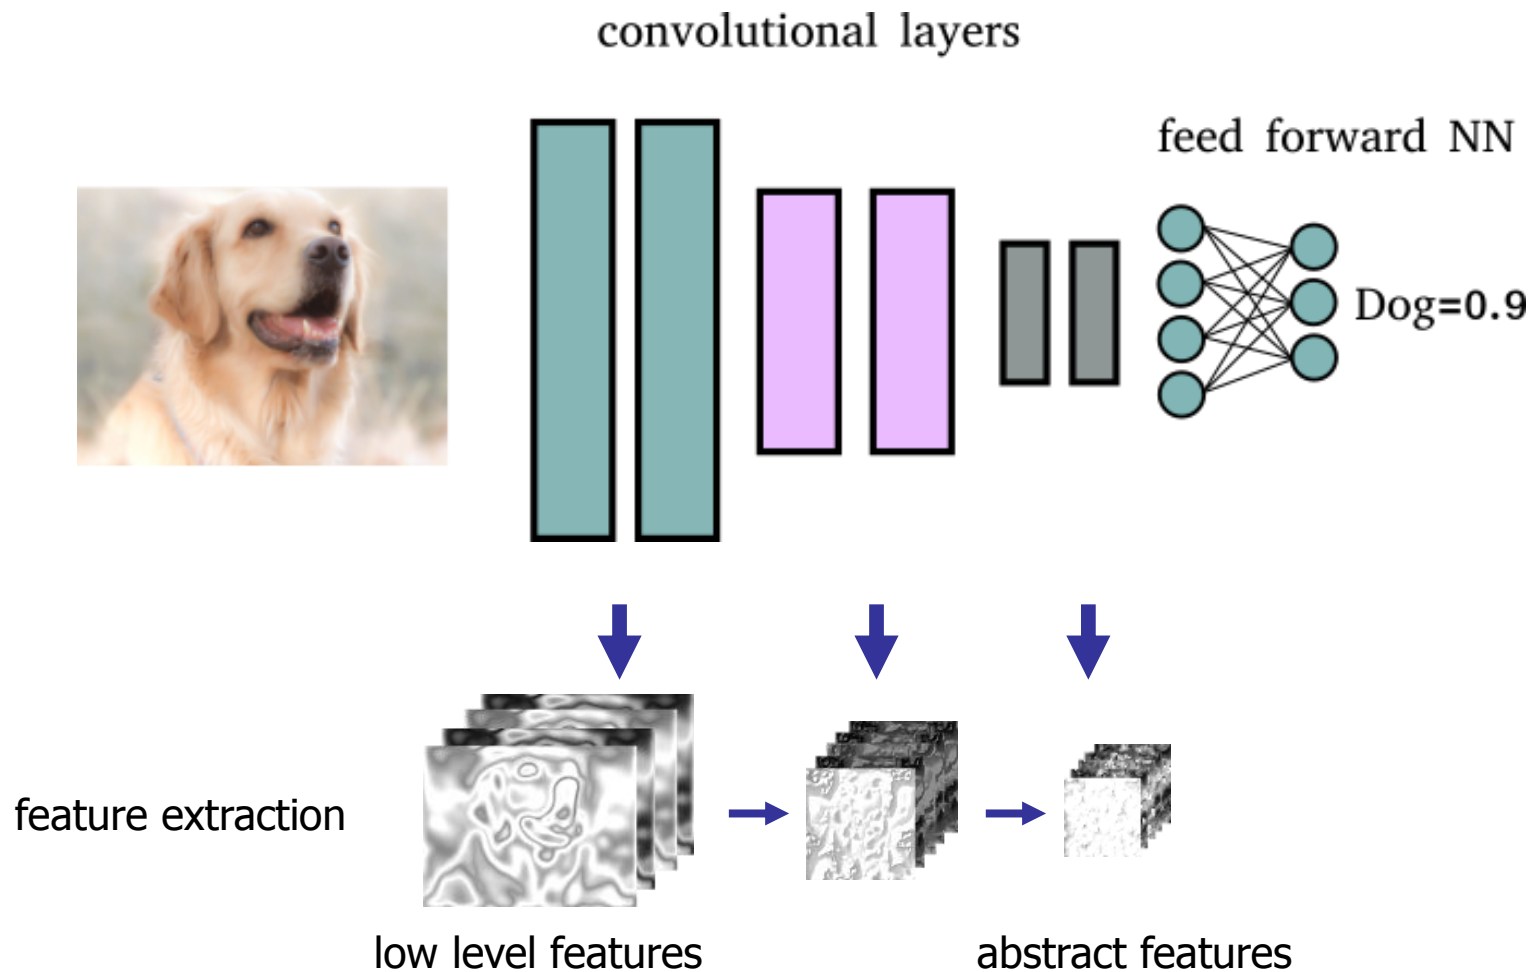
- Allow automatically extracting **features** from images and performing **classification**



Convolutional Neural Network (CNN) Architecture

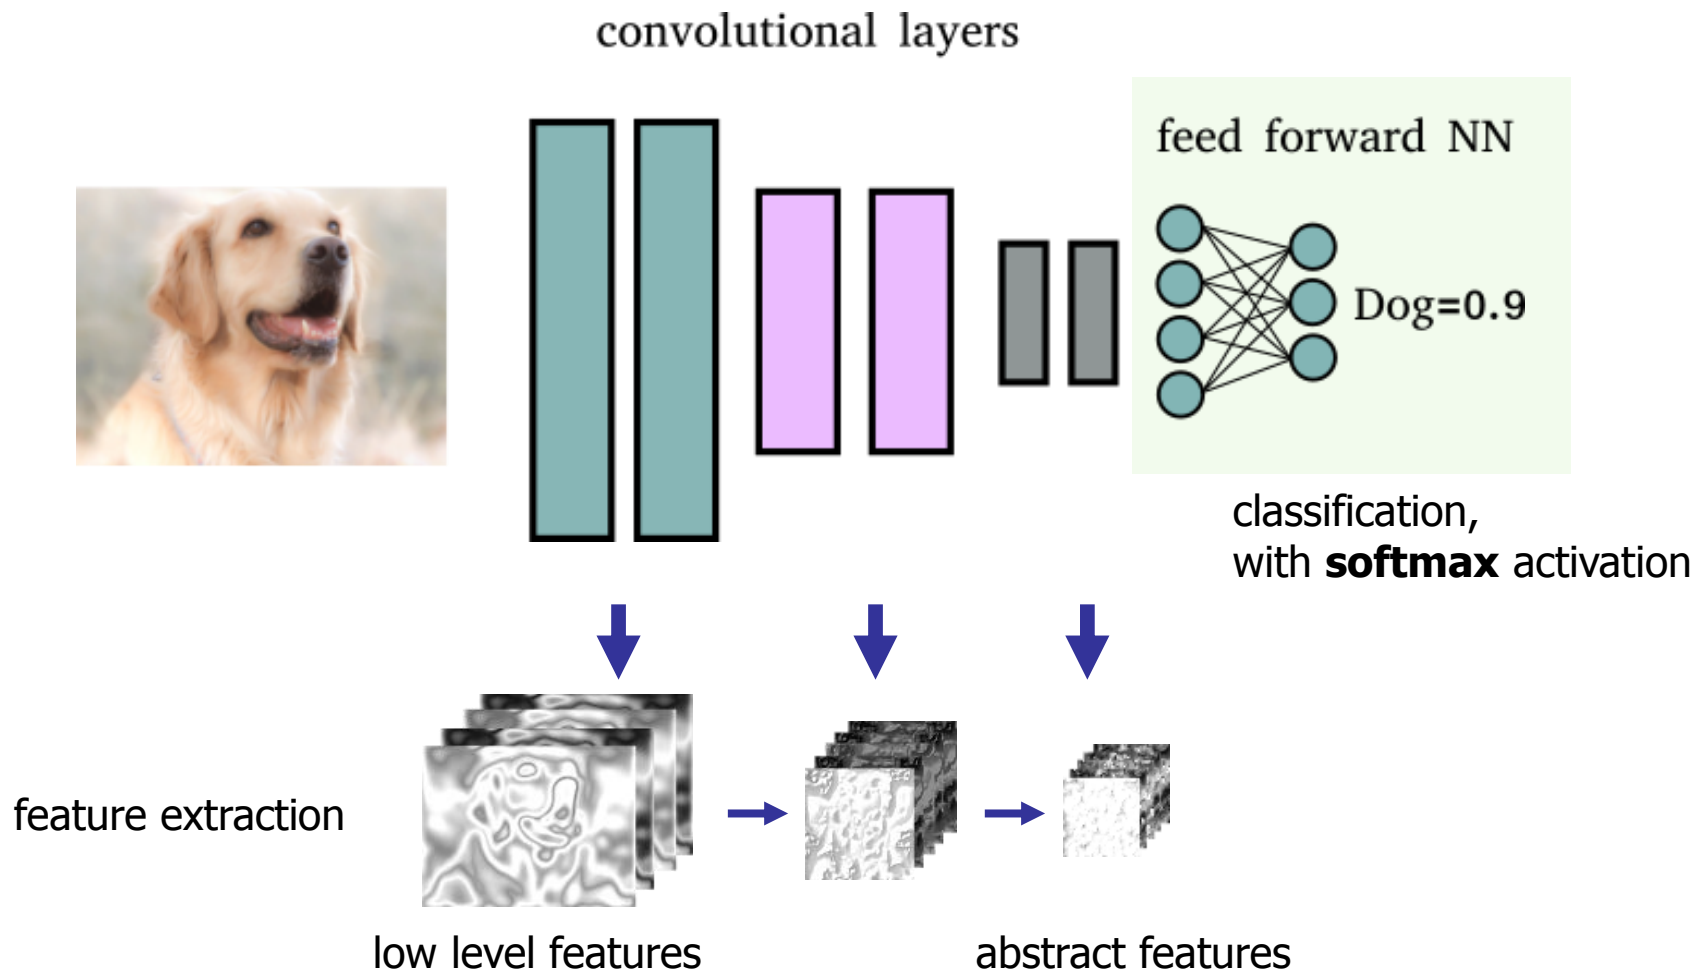


Convolutional Neural Networks





Convolutional Neural Networks






Convolutional Neural Networks

■ Typical convolutional layer

- *convolution* stage: feature extraction by means of (hundreds to thousands) sliding filters
- sliding filters *activation*: apply activation functions to input tensor
- *pooling*: tensor downsampling






Convolutional Neural Networks

■ Tensors

- data flowing through CNN layers is represented in the form of *tensors*
- *Tensor* = N-dimensional vector
- *Rank* = number of dimensions
 - scalar: rank 0
 - 1-D vector: rank 1
 - 2-D matrix: rank 2
- *Shape* = number of elements for each dimension
 - e.g. a vector of length 5 has shape [5]
 - e.g. a matrix $w \times h$, $w=5$, $h=3$ has shape $[h, w] = [3, 5]$

rank-3 tensor with shape
 $[d, h, w] = [4, 2, 3]$






Convolutional Neural Networks

■ Images

- rank-3 tensors with shape $[d, h, w]$
- where h =height, w =width, d =image depth (1 for grayscale, 3 for RGB colors)






Convolutional Neural Networks

■ Convolution

- processes data in form of *tensors* (multi-dimensional matrixes)
- **input**: input image or intermediate features (tensor)
- **output**: a tensor with the extracted features






Convolutional Neural Networks

■ Convolution

- a *sliding filter* produces the values of the output tensor
- sliding filters contain the trainable *weights* of the neural network
- each convolutional layer contains many (hundreds) filters






Convolutional Neural Networks

■ Convolution

- images are transformed into features by convolutional filters
- after convolving a tensor $[d, h, w]$ with N filters we obtain
 - a rank-3 tensor with shape $[N, h, w]$
 - hence, each filter generates a layer in the depth of the output tensor







Convolutional Neural Networks

■ Activation

- simulates biological activation to input stimuli
- provides non-linearity to the computation
- ReLU is typically used for CNNs
 - faster training (no vanishing gradients)
 - does not saturate
 - faster computation of derivatives for backpropagation






Convolutional Neural Networks

■ Pooling

- performs tensor *downsampling*
- *sliding filter* which replaces tensor values with a *summary* statistic of the nearby outputs
- *maxpool* is the most common: computes the maximum value as statistic






Convolutional Neural Networks

■ Convolutional layers training

- during training each sliding filter learns to recognize a particular *pattern* in the input tensor
- filters in *shallow layers* recognize textures and edges
- filters in *deeper layers* can recognize objects and parts (e.g. eye, ear or even faces)

shallow filters



deeper filters

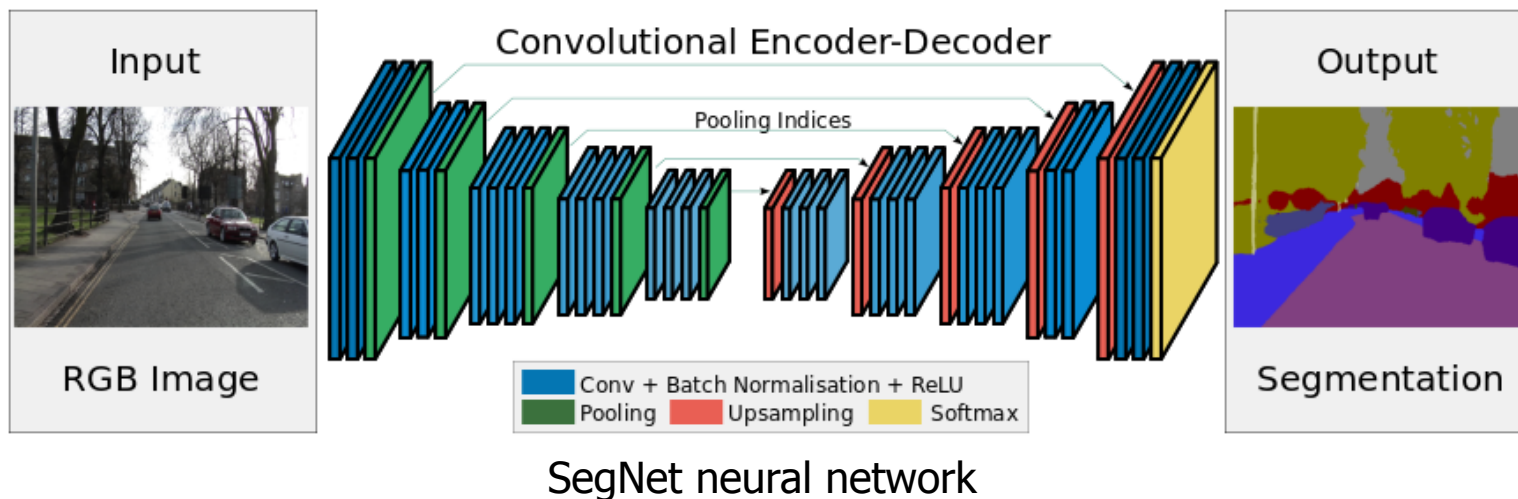




Convolutional Neural Networks

■ Semantic segmentation CNNs


- allow assigning a class to each pixel of the input image
- composed of 2 parts
 - **encoder network:** convolutional layers to extract abstract features
 - **decoder network:** deconvolutional layers to obtain the output image from the extracted features





Recurrent Neural Networks

- Allow processing *sequential* data $x(t)$
- Differently from normal FFNN they are able to keep a *state* which evolves during time
- Applications
 - machine translation
 - time series prediction
 - speech recognition
 - part of speech (POS) tagging






Recurrent Neural Networks

- RNN execution during time


instance of the RNN at time t_1





Recurrent Neural Networks

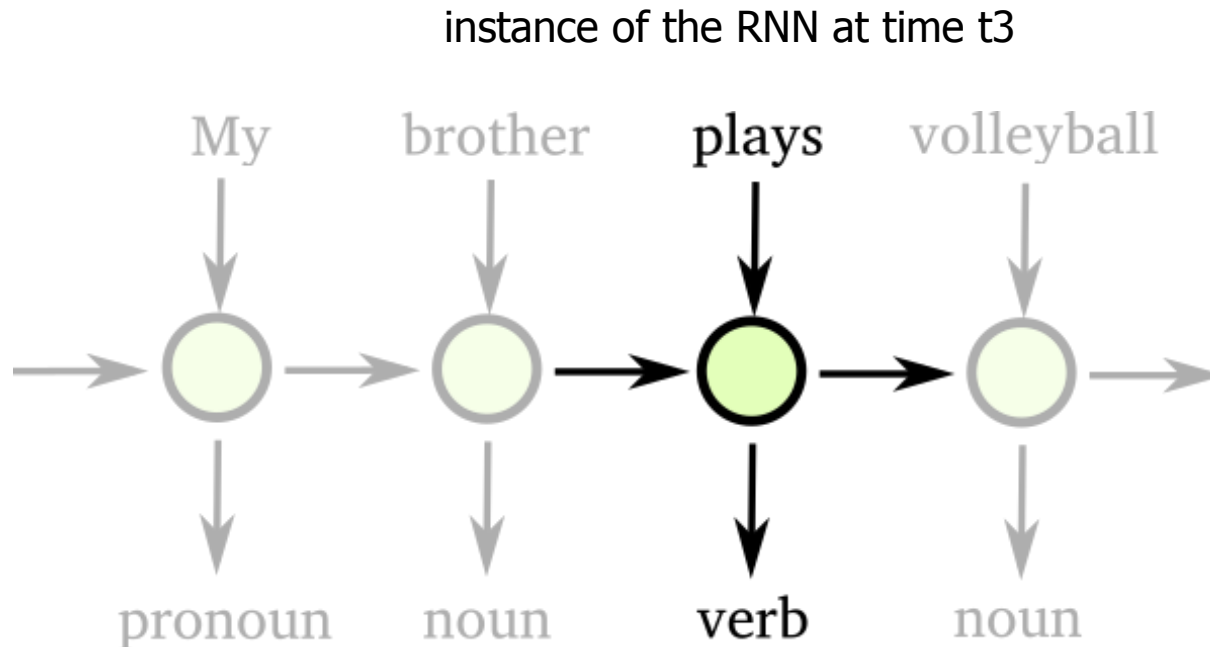
- RNN execution during time





Recurrent Neural Networks

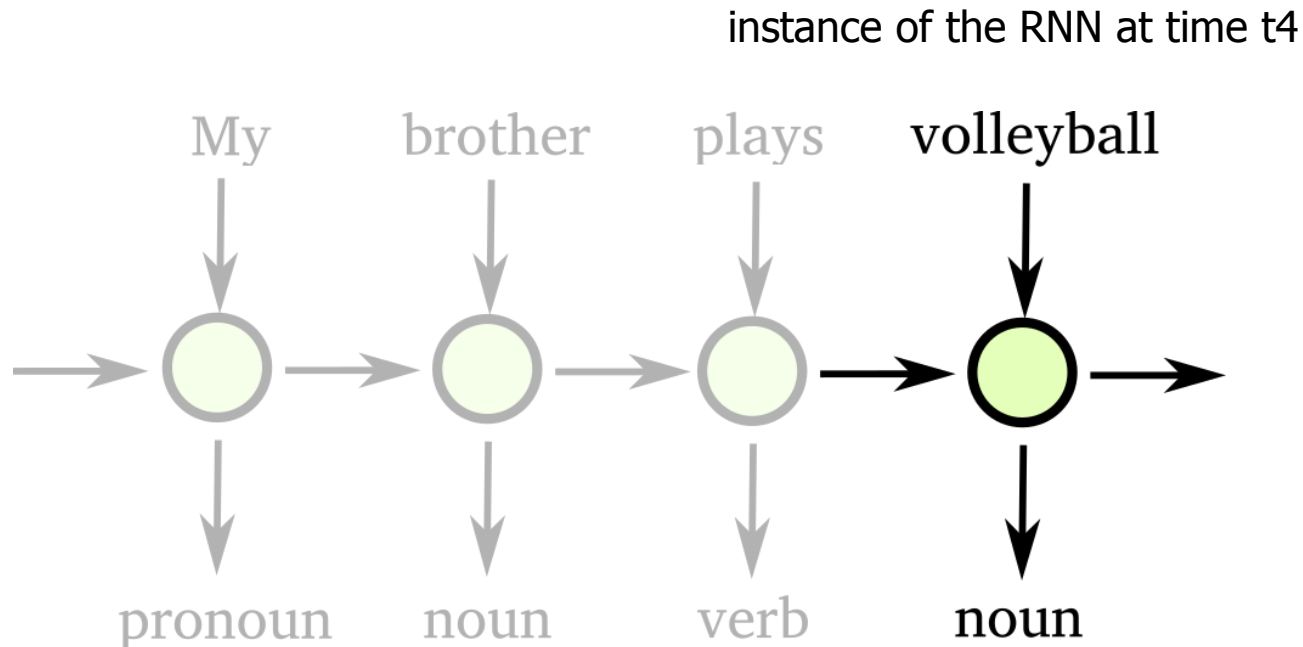
- RNN execution during time





Recurrent Neural Networks


- RNN execution during time





Recurrent Neural Networks

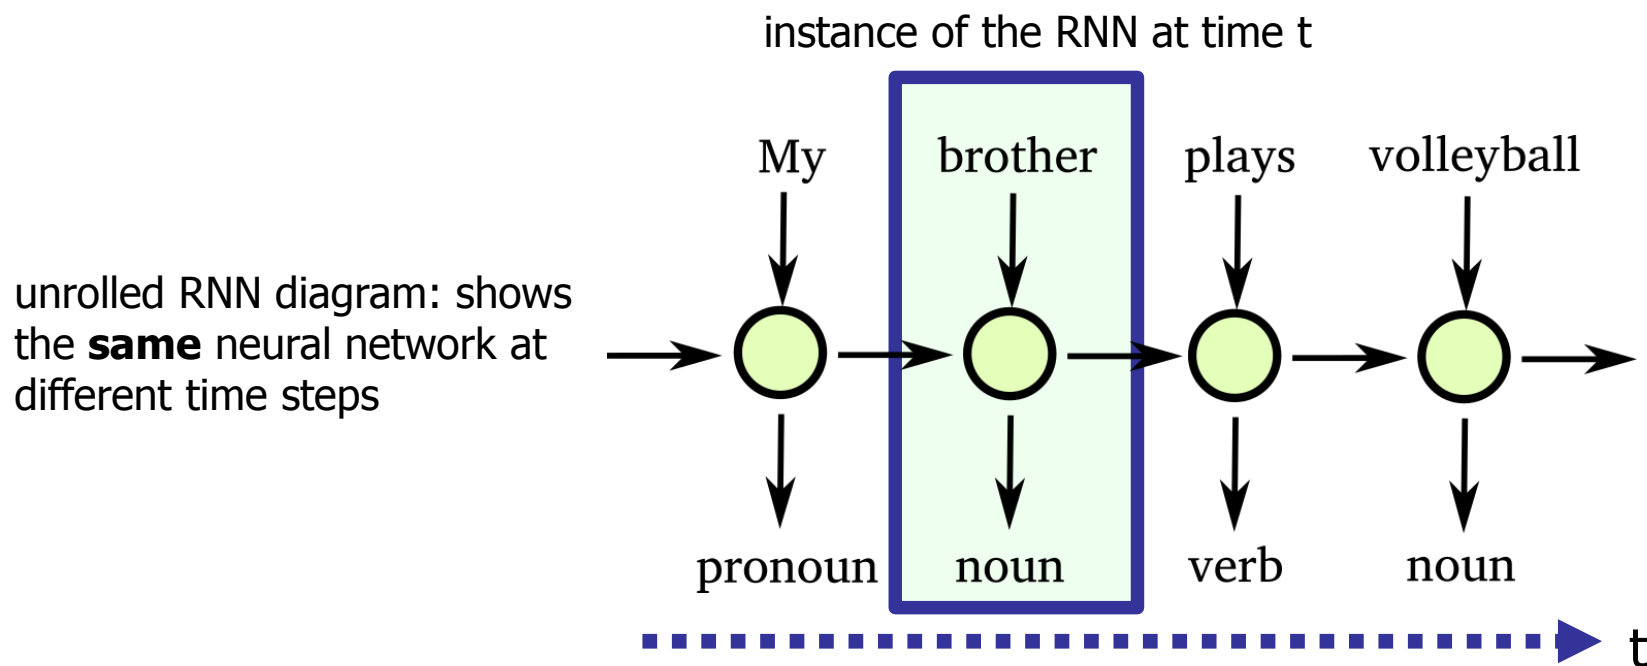
- A RNN receives as input a vector $x(t)$ and the state at previous time step $s(t-1)$
- A RNN typically contains many *neurons organized in different layers*





Recurrent Neural Networks

- Training is performed with *Backpropagation Through Time*
- Given a pair training sequence $x(t)$ and expected output $y(t)$
 - error is propagated through time
 - weights are updated to minimize the error across all the time steps






Recurrent Neural Networks

- Issues

- *vanishing gradient*: error gradient decreases rapidly over time, weights are not properly updated
- this makes harder having RNN with *long-term* memories

- Solution: *LSTM* (Long Short Term Memories)

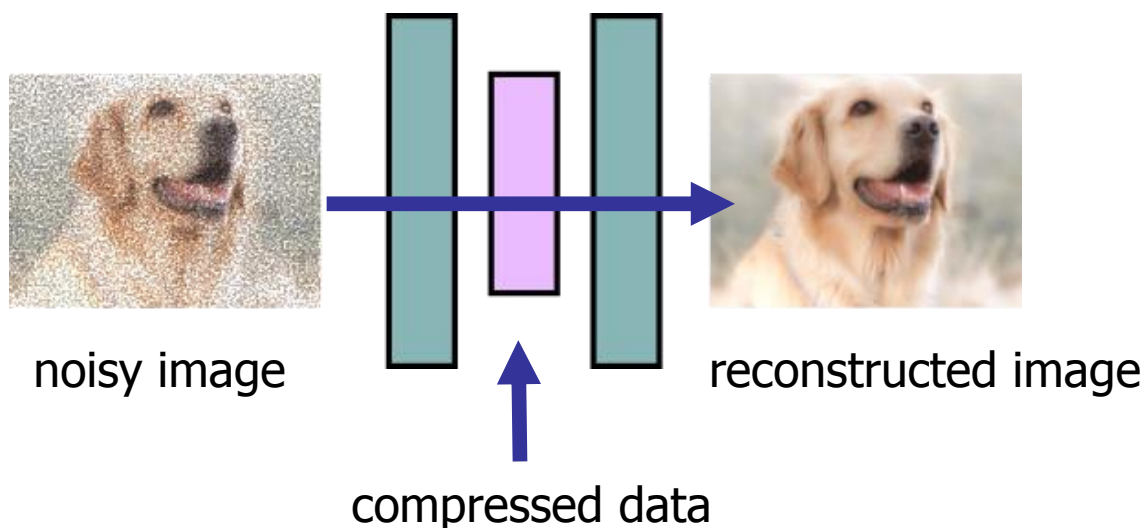
- RNN with “gates” which encourage the state information to flow through long time intervals





Autoencoders

- Autoencoders allow *compressing* input data by means of compact representations and from them *reconstruct* the initial input
 - for feature extraction: the compressed representation can be used as significant set of features representing input data
 - for image (or signal) *denoising*: the image reconstructed from the abstract representation is denoised with respect to the original one







Word Embeddings (Word2Vec)

- Word *embeddings* associate words to n-dimensional vectors
 - trained on big text collections to model the word distributions in different sentences and contexts
 - able to capture the *semantic* information of each word
 - words with similar *meaning* share vectors with similar characteristics




input word





Word Embeddings (Word2Vec)



- Since each word is represented with a vector, operations among words (e.g. difference, addition) are allowed





Word Embeddings (Word2Vec)

- Semantic relationships among words are captured by vector positions



$\text{king} - \text{man} = \text{queen} - \text{woman}$
 $\text{king} - \text{man} + \text{woman} = \text{queen}$