

ANALYSIS OF BIG DATA ARCHITECTURES

Dimensions of Big Data architectures

- **Data model(s):**
 - Relations, trees (XML, JSON), graphs (RDF, others...), nested relations
 - Query language
- **Heterogeneity** (DM, QL): none, some, a lot
- **Scale:** small (~10-20 sites) or large (~10.000 sites)
- **ACID** properties
- **Control:**
 - Single master w/complete control over N slaves (Hadoop/HDFS)
 - Sites publish independently and process queries as directed by single master/*mediator*
 - Many-mediator systems, or peer-to-peer (P2P) with *super-peers*
 - Sites completely independent (P2P)

Architectures we will cover

- Distributed databases
- Mediator (data integration) systems
- Dataspaces, data lakes
- Peer-to-peer data management systems
- Structured data management on top of MapReduce


DISTRIBUTED RELATIONAL DATABASES

Distributed relational databases


- Oldest distributed architecture ('70s): IBM System R*
- Illustrate/introduce the main principles
- **Data** is distributed among many *nodes* (*sites, peers...*)
 - **Data catalog**: information on which data is stored where
 - Explicit : « All Paris sales are stored in Paris ».
Horizontal/vertical table fragmentation
Catalog stored at a master/central server.
 - Implicit: « Data is distributed by the value of the city »
(« somewhere »)
- **Queries** are distributed (may come from any site)
- **Query processing** is distributed
 - Operators may run on different sites → network transfer
 - Another layer of complexity to the optimization process

Distributed query optimization

Example 1: $R@s1$, $S@s2$, $T@s3$, $q@s4$



Example 2: $R@s1$, $S@s2$, $T@s3$, $U@s4$, $q@s5$




Plan pruning criteria if all the sites and network connections have equal performance:


- Ship the smaller collection

Distributed query optimization

Example 1: $R@s1$, $S@s2$, $T@s3$, $q@s4$



Example 2: $R@s1$, $S@s2$, $T@s3$, $U@s4$, $q@s5$




Plan pruning criteria if all the sites and network connections have equal performance:


- Ship the smaller collection
- Transfer to join partner or the query site

Distributed query optimization

Example 1: $R@s1$, $S@s2$, $T@s3$, $q@s4$



Example 2: $R@s1$, $S@s2$, $T@s3$, $U@s4$, $q@s5$




Plan pruning criteria if all the sites and network connections have equal performance:

- Ship the smaller collection.
- Transfer to join partner or the query site

This plan illustrates total effort != response time

Distributed query optimization technique: semijoin reducers

- $R \text{ join } S = (R \text{ semijoin } S) \text{ join } S$



- Useful in distributed settings to reduce transfers: *if the distinct $S.b$ values are smaller than the non-matching R tuples*
- Symmetrical alternative: $R \text{ join } S = R \text{ join } (S \text{ semijoin } R)$
- This gives one more alternative in every join \rightarrow search space explosion
- Heuristics [Stocker, Kossmann et al., ICDE 2001]

Distribution of control in distributed relational DBs (1970s)

Servers DB1@site1: R1(a,b), S1(a,c)

Server DB2@site2: R2(a,b), S2(a,c),

Server DB3@site3: R3(a,b),
S3(a,c) defined as:

```
select * from DB1.S1 union all
select * from DB2.S2 union all
select R1.a as a, R2.b as c from DB1.R1 r1, DB2.R2 r2
where r1.a=r2.a
```

DB3@site3 decides what to import from site1, site2 (« hard links »)

Site1, site2 are independent servers

Also: replication policies, distribution etc. (usually with one or a few masters)

Modern distributed databases: H-Store (subsequently VoltDB), 2016

- From the team of Michael Stonebraker (Turing Award, author of the Postgres system)
 - H-Store: research prototype
 - VoltDB: commercial product issued from H-Store
- Main goal: quick OLTP (**online transaction processing**), e.g., sales, likes, posts...
- Built to run on **cluster** for horizontal scalability
- **Share-nothing architecture**: each node stores tables **shards** (+ k replication for durability)

Frequent concept in Big Data architectures: shards

- **Shard** = small fragment of a data collection (e.g., a table)
- The assignment of data items (e.g. tuples) into shards is often done by **hashing** on tuple key
 - The table must have at least one key
 - Hashing ensures (with high probability) uniform distribution
- Key-based hashing is used as a mechanism for implementing distributed data catalogs. We will encounter it often.



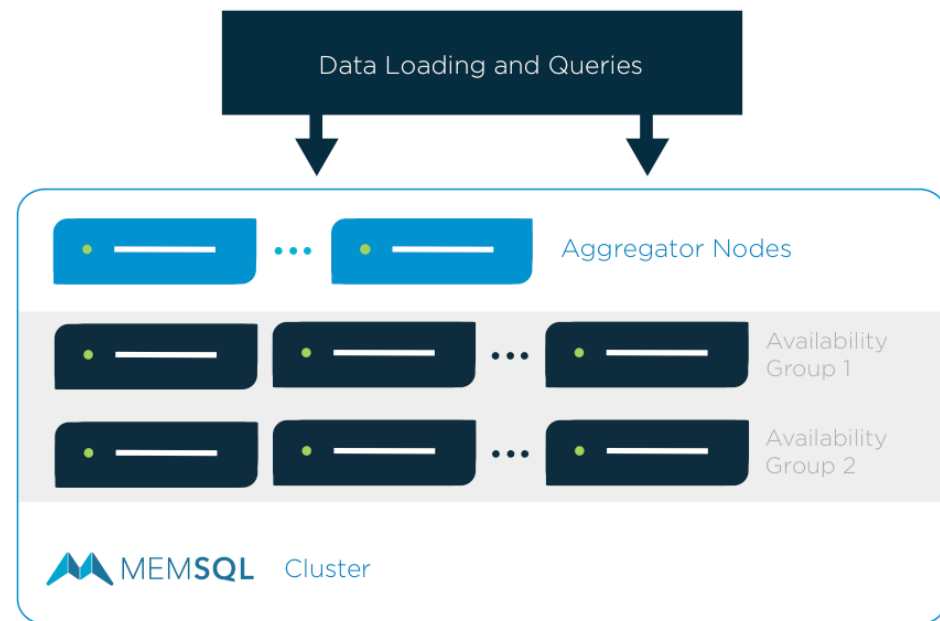
H-Store transactions

- Applications call **stored procedures** = code which also contains SQL queries
 - Each contained SQL query is partially unknown (depends on parameters specified at runtime); H-Store "pre-optimizes" it
- 1 **transaction** = 1 call of a stored procedure
- Can be submitted to any node, together with parameters
- The node can run the procedure up to the query(ies) → updated, completely known plan → transaction manager

Modern distributed RDB: MemSQL (2013)

MemSQL runs with

- a **master aggregator**, responsible of the metadata (catalog)
- possibly more aggregators
- at least one **leaf**, each of which stores part(s) of some table(s)
- In each leaf, there are **partitions** (by default: 1 per CPU core)



Availability group: a set of machines + a set of replica machines (one-to-one)

Query processing in MemSQL

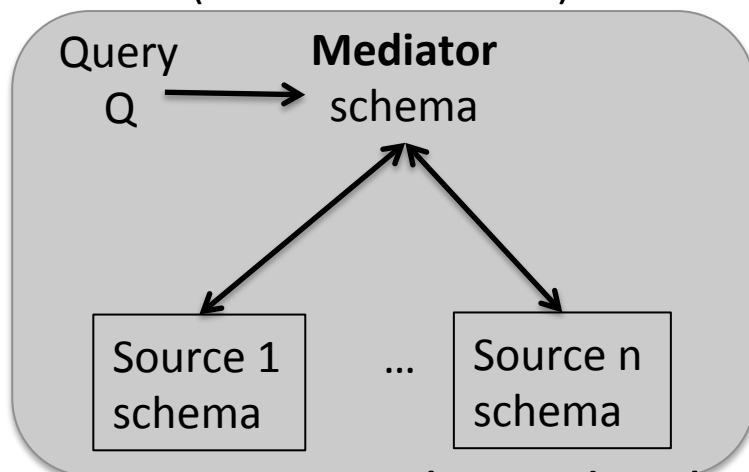
- **Indexes** managed within each partition
- In general, every query is run with a level of parallelism equal to the number of partitions
- **Select** queries are executed by the leaves which hold some partition(s) with data matching the query
- **Aggregation** queries run at the leaves involved and at the aggregator(s)
- **Join** queries
 - Easy if one input is a *reference* (small) table: one that is replicated fully to every machine in the cluster
 - Otherwise, they recommend **sharing the shard key across tables to be joined**
 - Also called **co-partitioning**, we will be seeing this again
 - Otherwise, joins will incur data transfer within the cluster.

MEDIATOR SYSTEMS

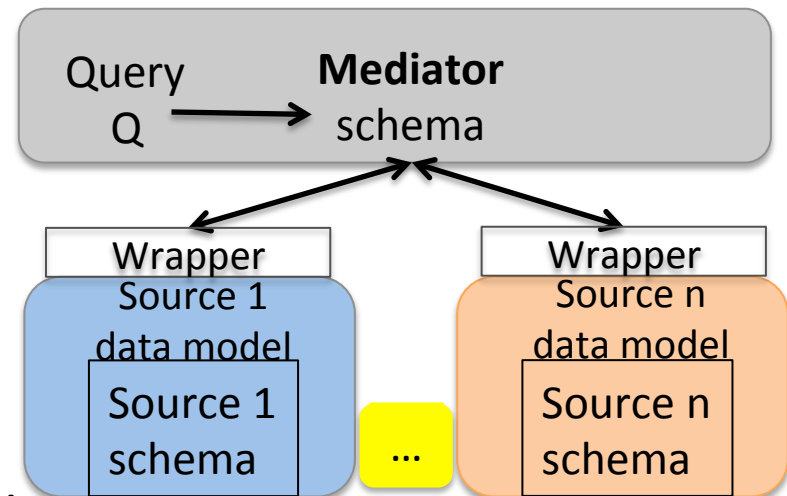
Mediator systems

- A set of **data sources**, each with: data model, query language, and schema (also called source schemas).
 - DM and QL may differ across sources
- A **mediator** with its own DM, QL and mediator schema
- Queries are asked against the mediator schema

Common data model
(sources+mediator)



Mediator data model



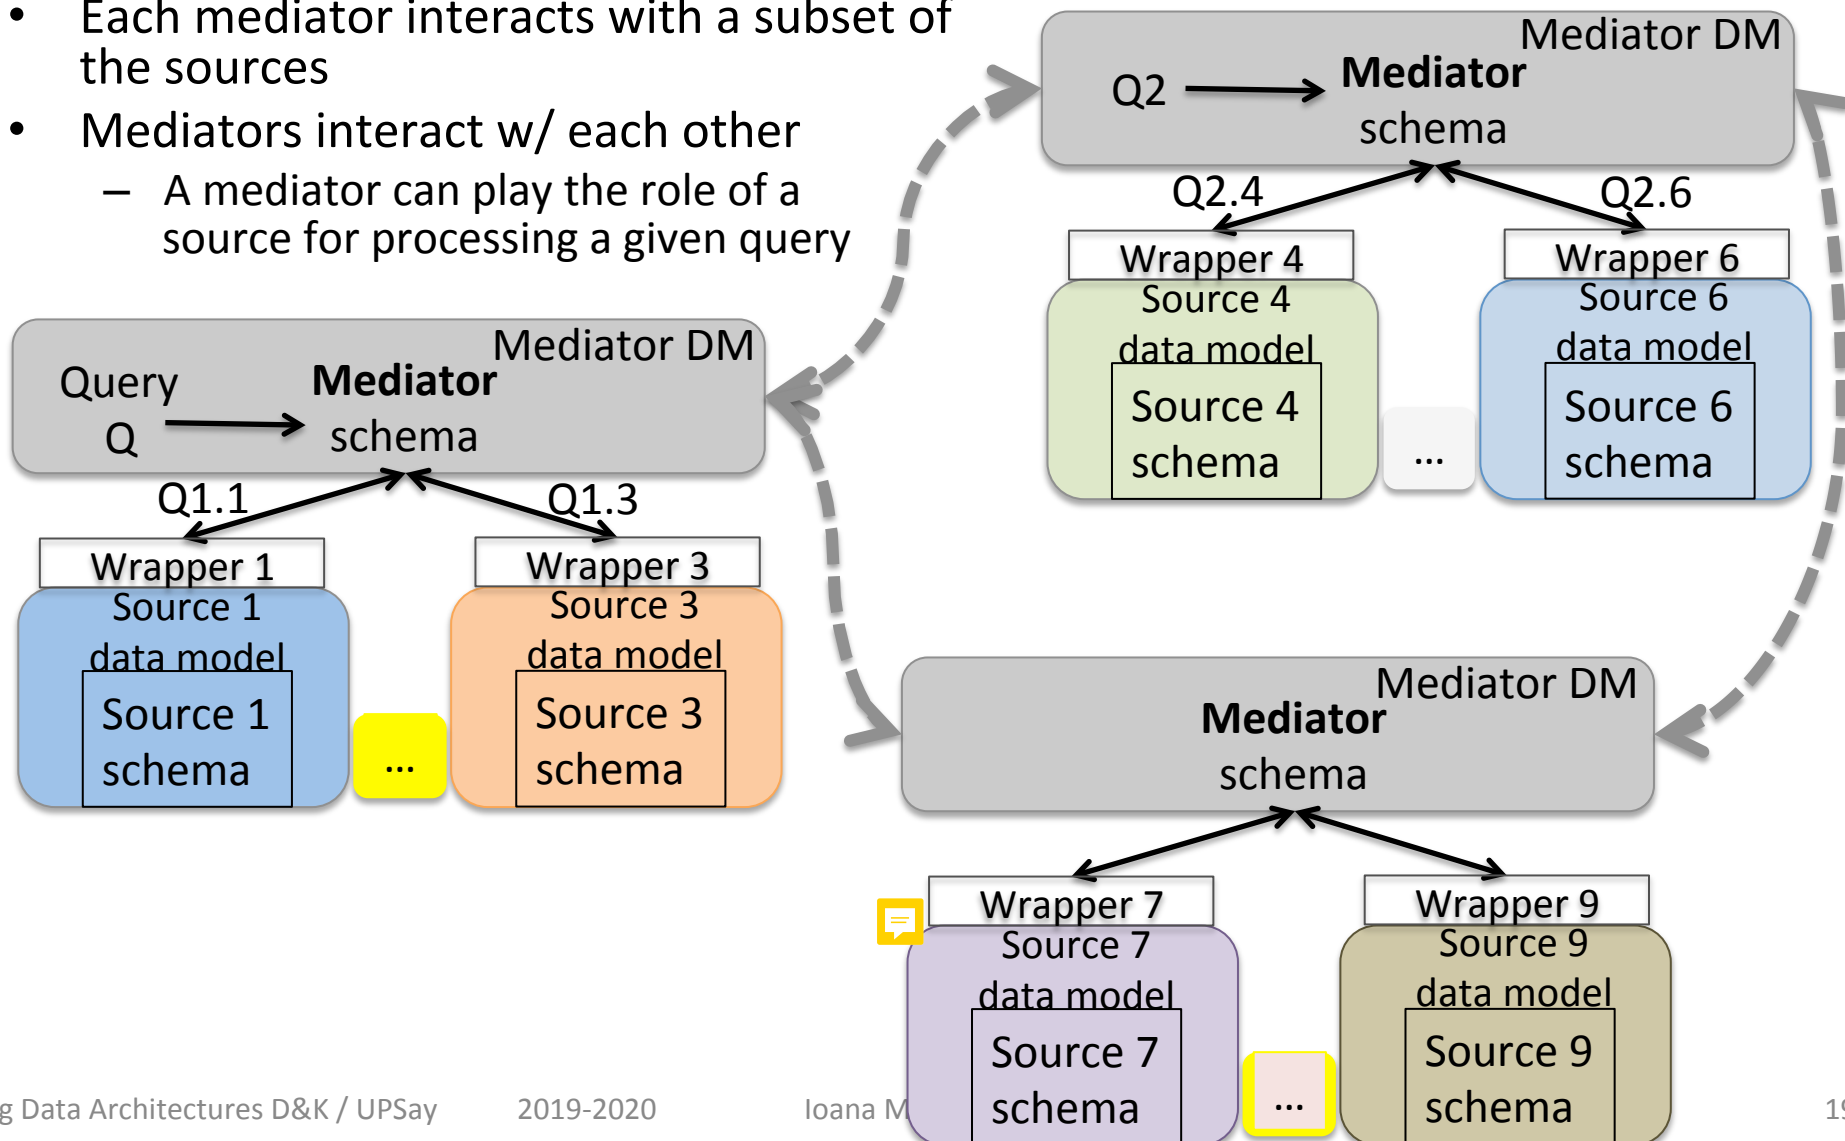
- **ACID**: mostly read-only; **size**: small
- **Control**: Independent publishing; mediator-driven integration

EXAMPLE: TATOOINE DEMO

[VLDB 2016, [HTTPS://TEAM.INRIA.FR/CEDAR/TATOOINE/](https://team.inria.fr/cedar/tatooine/)]

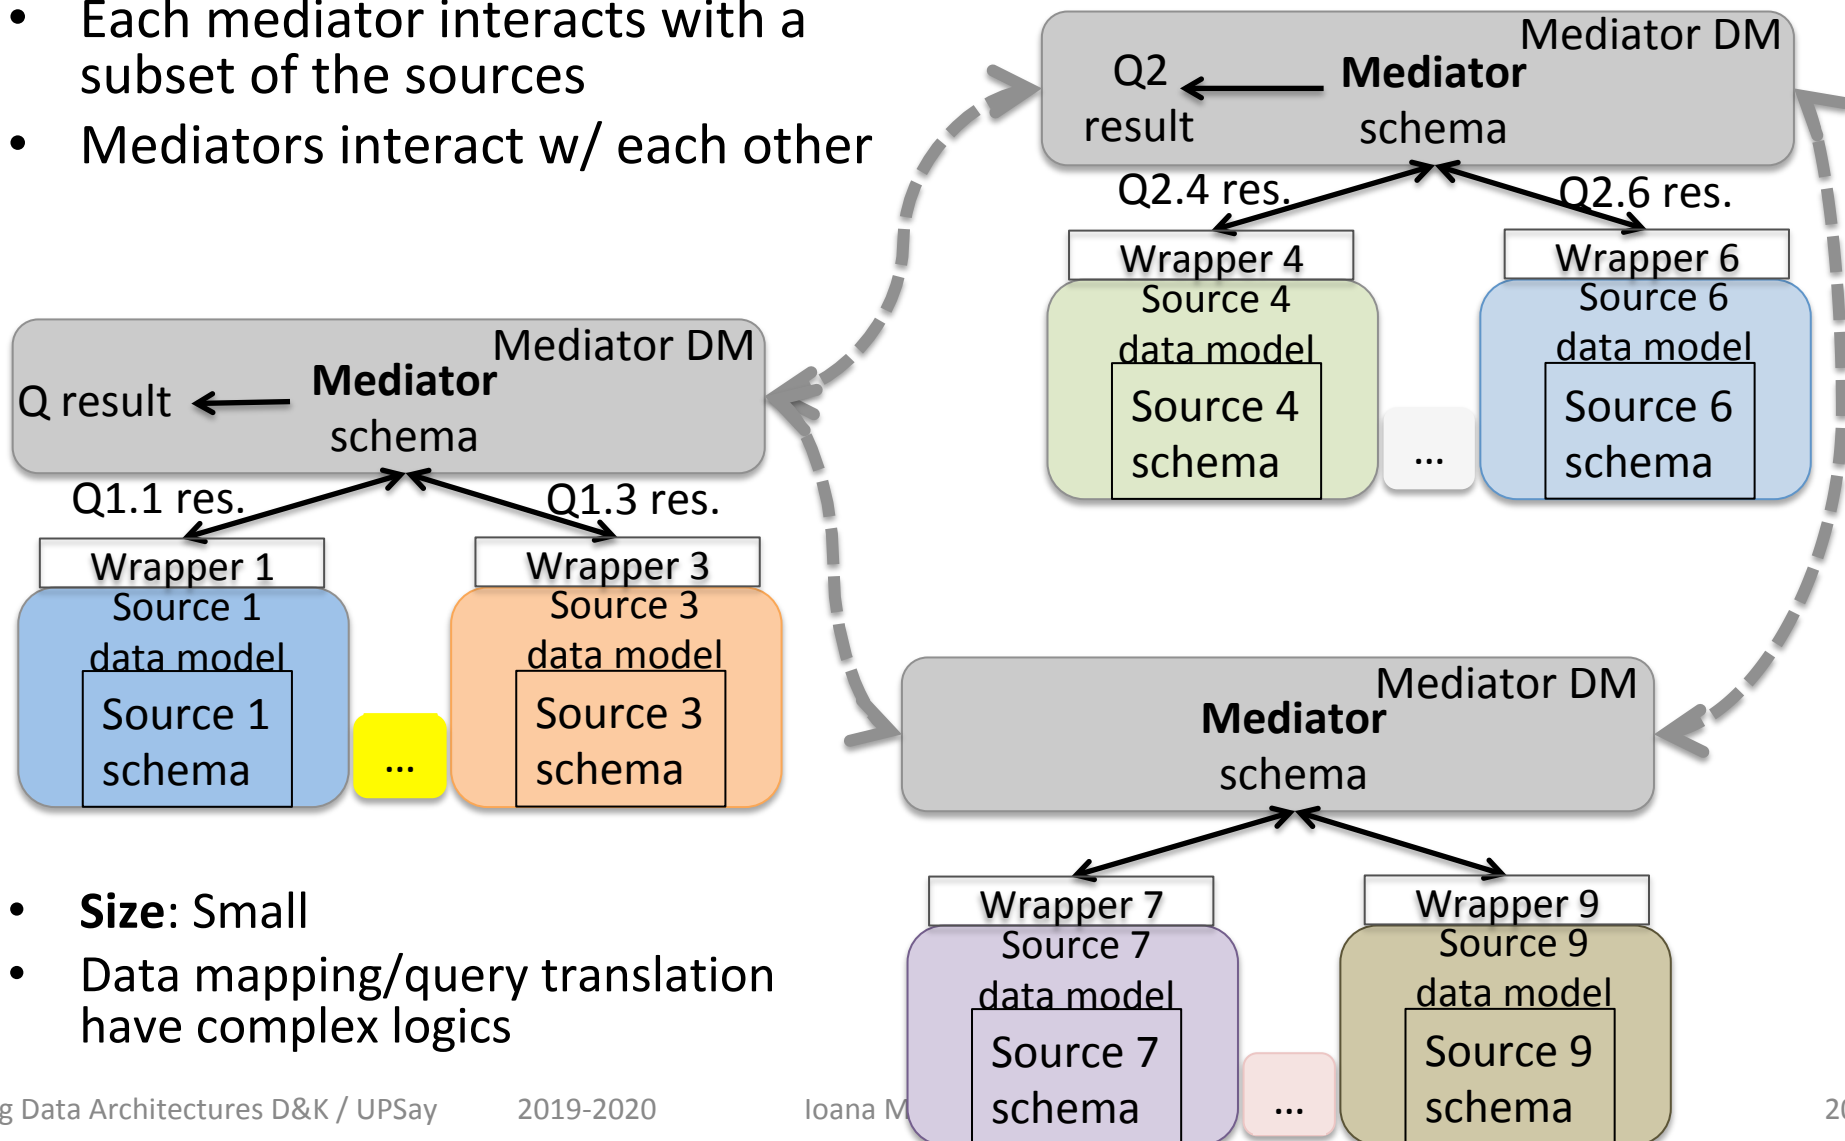
Many-mediator systems

- Each mediator interacts with a subset of the sources
- Mediators interact w/ each other
 - A mediator can play the role of a source for processing a given query



Many-mediator systems

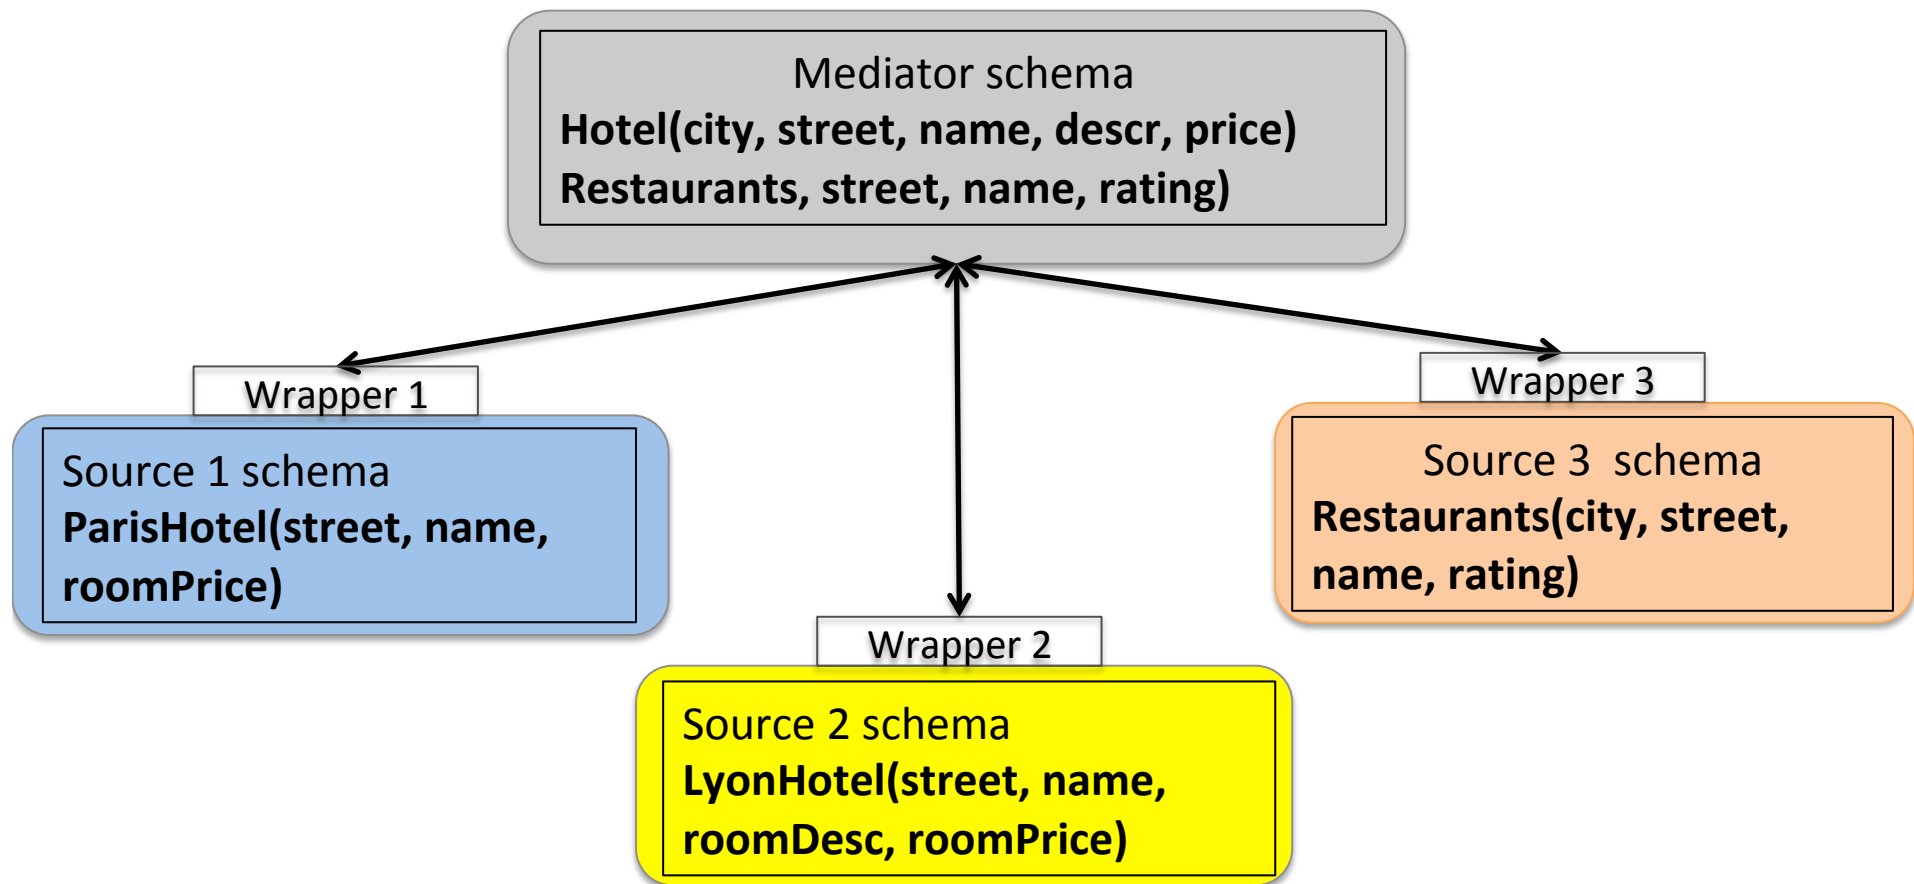
- Each mediator interacts with a subset of the sources
- Mediators interact w/ each other



- **Size:** Small
- Data mapping/query translation have complex logics

Connecting the source schemas to the global schema

- Sample scenario:



Connecting the source schemas to the global schema

- Data only exists in the sources.
- Applications only have access to, and only query, the mediator schema.
- How to express the **relation** between
 - the **mediator schema** accessible to applications, and
 - the **source schemas** reflecting the real data
 - so that a query over the mediator schema can be **automatically translated** into a query over the source schemas?
- Three approaches exist (see next)

Connecting the source schemas to the global schema: Global-as-view (GAV)

s1:ParisHotels(street, name, roomPrice)

s2:LyonHotel(street, name, roomDesc, roomPrice)

s3:Restaurant(city, street, name, rating)

Global: Hotel(city, street, name, descr, price),
Restaurant(city, street, name, rating)

Defining **Hotel** as a view over the source schemas:

define view Hotel as

select 'Paris' as city, street, name, null as roomDesc, roomPrice as price
from s1:ParisHotels

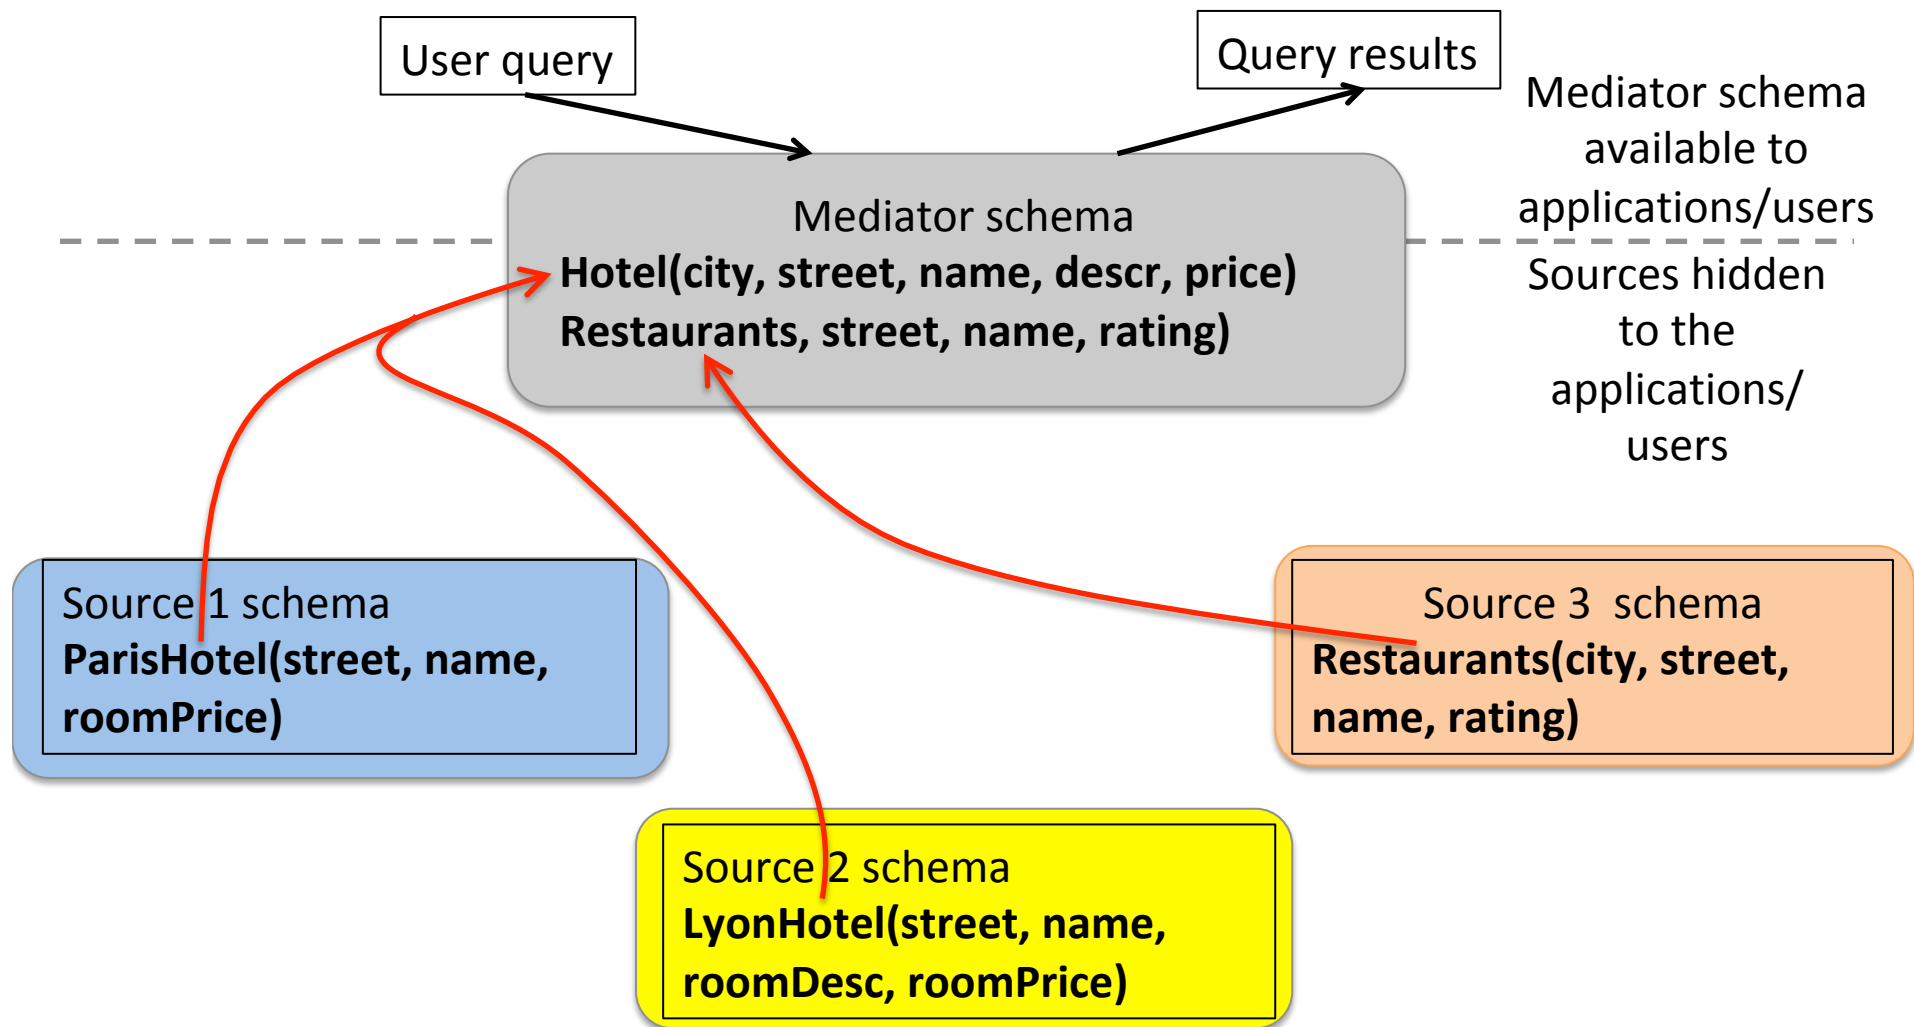
union all

select 'Lyon' as city, street, name, descr as roomDesc, price
from s2:LyonHotel

Defining **Restaurant** as a view over the source schemas:

define view Restaurant as select * from s3:Restaurant

Connecting the source schemas to the global schema: Global-as-View



Query processing in global-as-view (GAV)

```
define view Hotel as  
select 'Paris' as city, street, name, null as roomDesc, roomPrice as price  
from s1:ParisHotels  
union all  
select 'Lyon' as city, street, name, descr as roomDesc, price  
from s2:LyonHotel
```

Query:

```
select * from Hotel where city='Paris' and price<200      becomes:  
select * from (select 'Paris' as city... union... select 'Lyon' as city...)  
              where city='Paris' and price < 200          which becomes:  
select * from (select 'Paris' as city...)  
              where city='Paris' and price < 200          which becomes:  
select * from s1:ParisHotels where price < 200
```

Query processing in global-as-view (GAV)

```
define view Hotel as  
select 'Paris' as city, street, name, null as roomDesc, roomPrice as price  
from s1:ParisHotels  
  
union all  
  
select 'Lyon' as city, street, name, descr as roomDesc, price from s2:LyonHotel  
define view Restaurant as select * from s3:Restaurant
```

Query:

```
select h.street, r.rating from Hotel h, Restaurant r where h.city=r.city and  
r.city='Lyon' and h.street=r.street and h.price<200
```

becomes:

```
select h.street, r.rating from (select 'Paris' as city... from s1:ParisHotels  
union all select 'Lyon' as city... from s2:LyonHotel) h, (select * from s3:Restaurant) r  
where h.city=r.city and r.city='Lyon' and h.street=r.street and h.price<200
```

which becomes:

```
select h.street,r.rating from (select ... from s2:LyonHotel) h, s3:Restaurant r where  
r.city='Lyon' and h.street=r.street and h.price<200
```

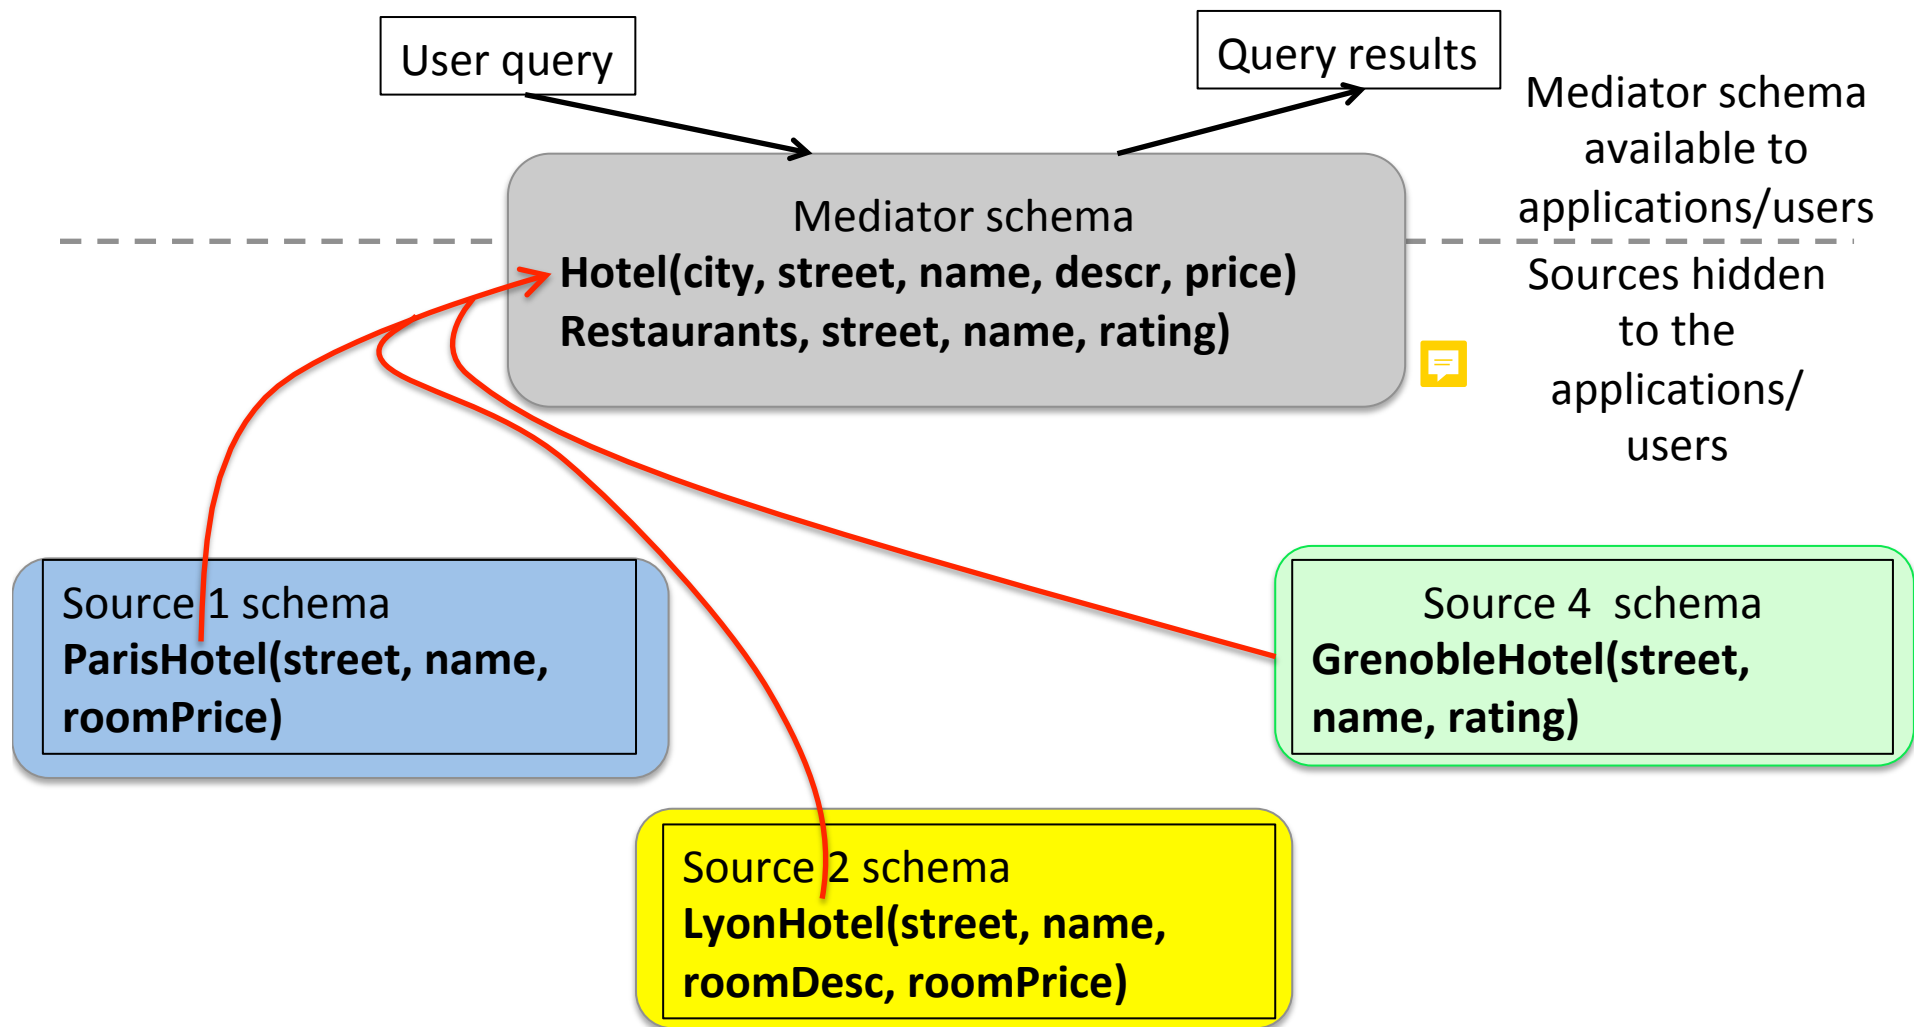
which becomes:

```
select h.street, r.rating from s2:LyonHotel h, s3:Restaurant r where r.city='Lyon' and  
h.price<200 and h.street=r.street
```

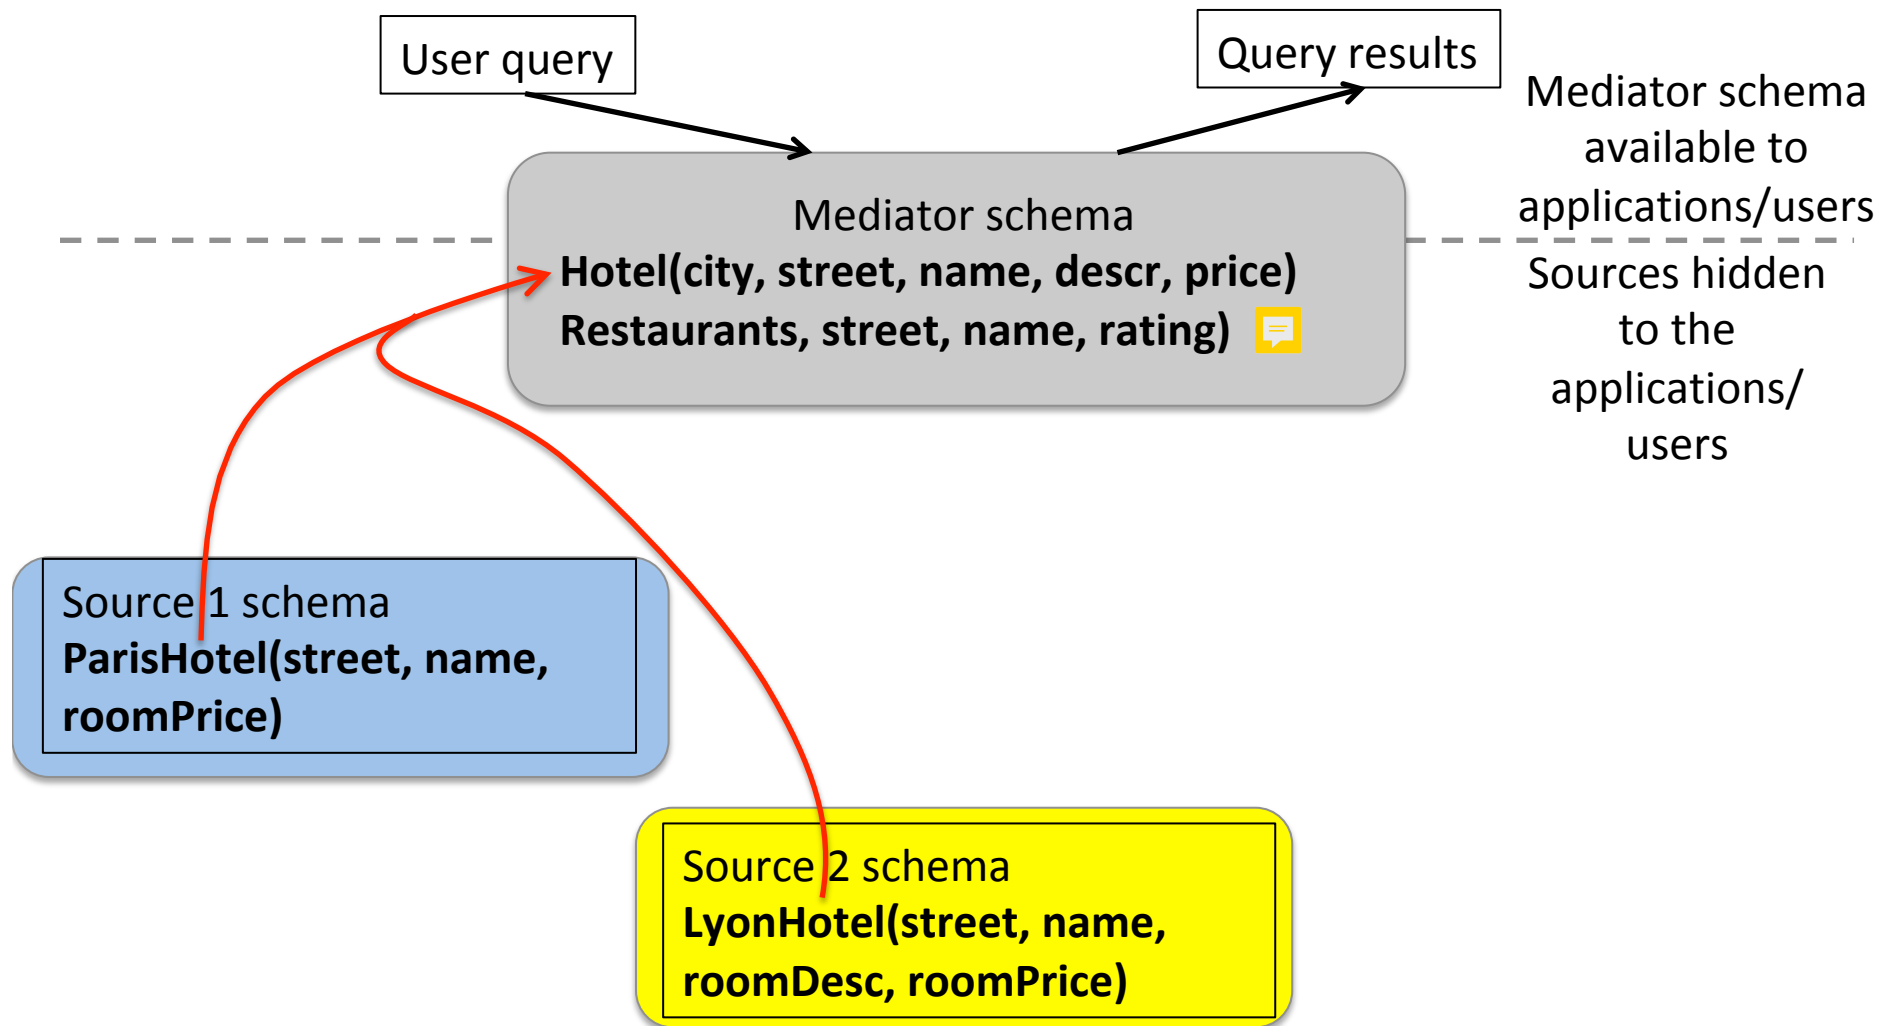
Concluding remarks on global-as-view (GAV)

- Query processing = **view unfolding**: replacing the view name with its definition and working out simple equivalences from there
 - Allows to push to each data source as much as it can do (trusted heuristic)
- **Weakness**: changes in the data sources require changes of the global schema
 - In the worst case, all applications written based on this global schema need to be updated
 - Hard to maintain

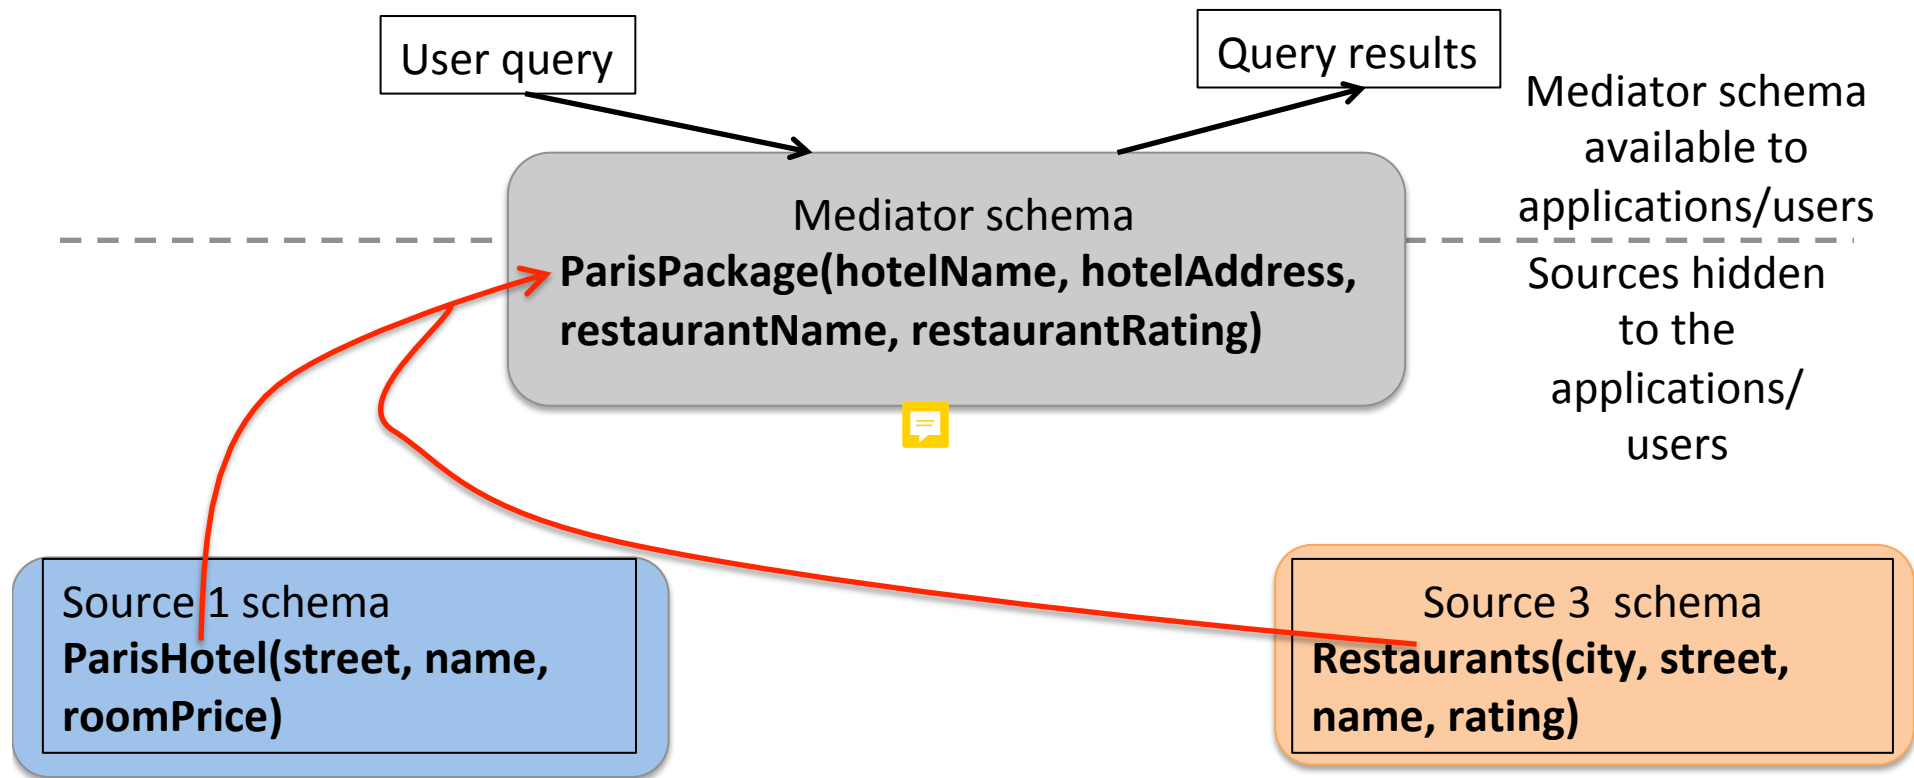
Global-as-View: Adding a new source



Global-as-View: Removing a source (1)



Global-as-View: removing a source (2)



If **Source3.Restaurant** withdraws, the **ParisPackage** relation in the global schema becomes empty; applications cannot even access **Source1.ParisHotels**, even though they are still available.

Connecting the source schemas to the global schema: Local-as-view (LAV)

s1:ParisHotel(street, name, roomPrice)

s2:LyonHotel(street, name, roomDesc, roomPrice)

s3:Restaurant(city, street, name, rating)

Global: Hotel(city, street, name, descr, price), **Restaurant**(city, street, name, rating)

Defining **s1:ParisHotels** as a view over the global schema:

```
define view s1:ParisHotels as  
select street, name, price as roomPrice  
from Hotel where city='Paris'
```

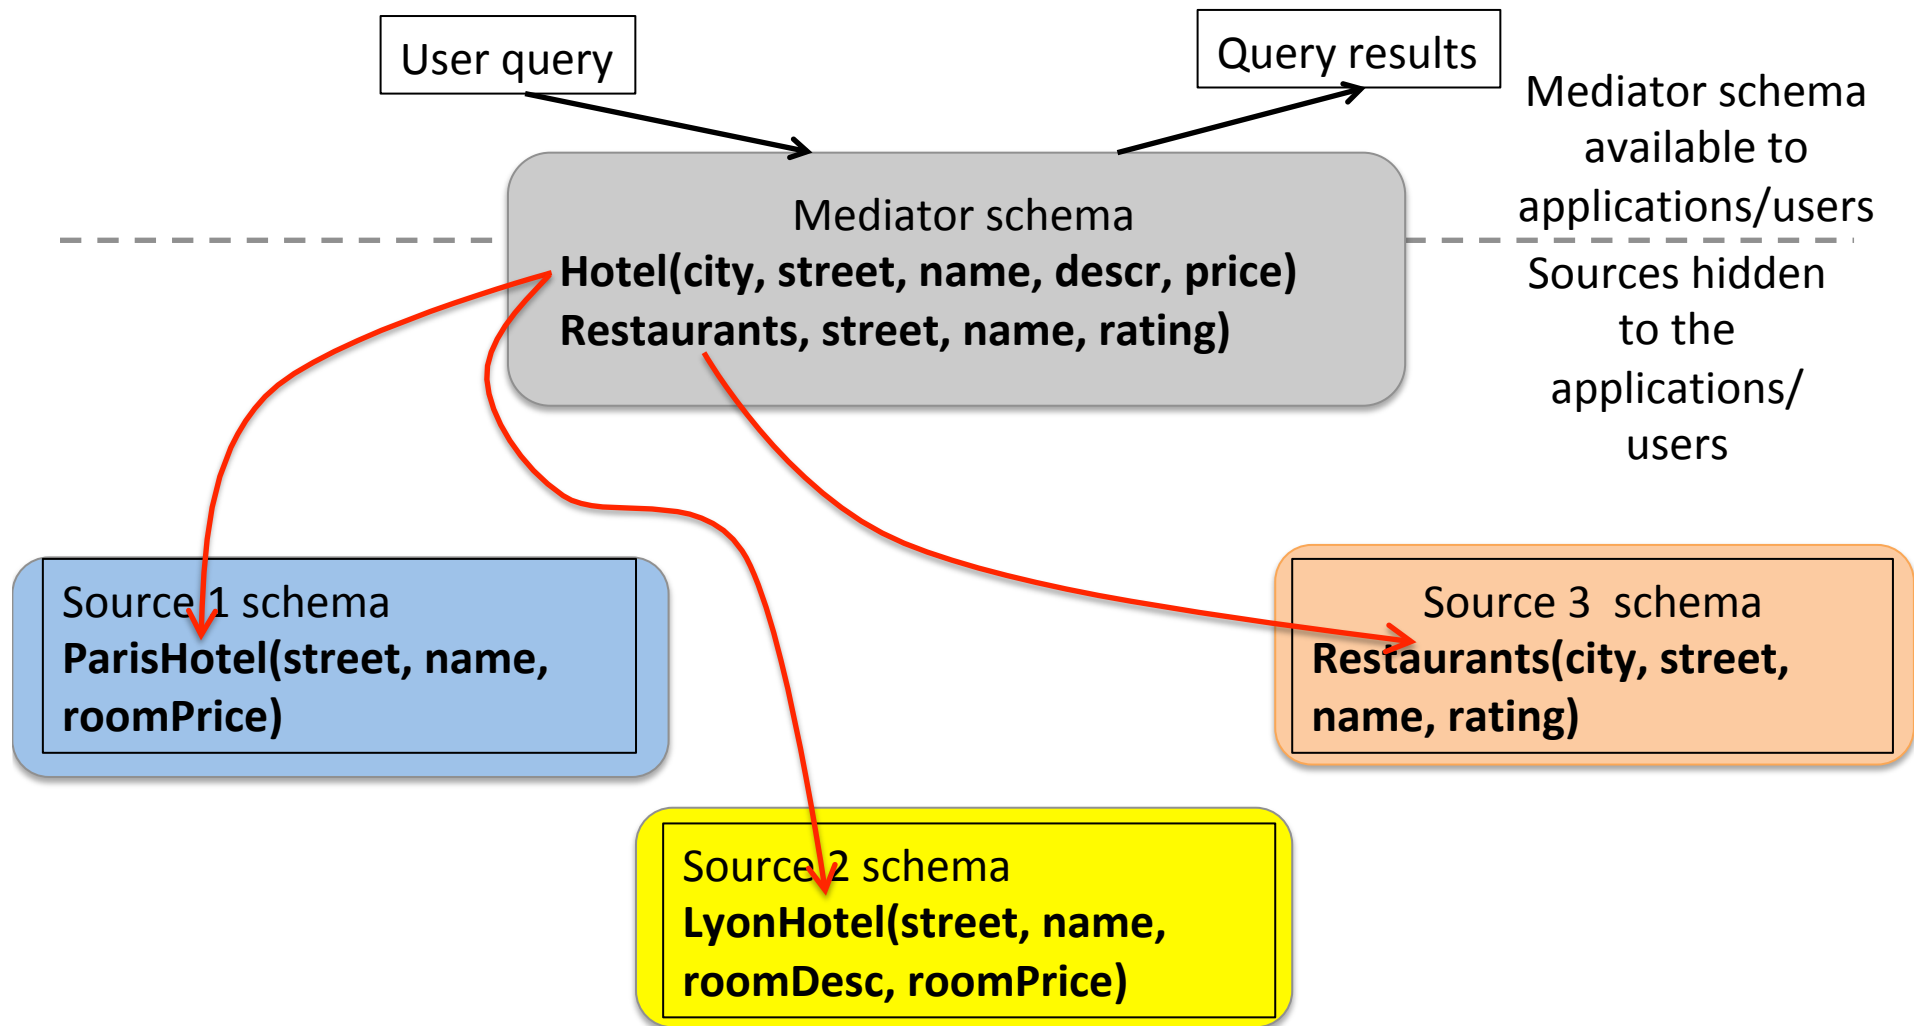
Defining **s2:LyonHotel** as a view over the global schema:

```
define view s2:LyonHotel as  
select street, name, descr as roomDesc, price as roomPrice  
from Hotel where city='Lyon'
```

Defining **s3:Restaurant** as a view over the global schema:

```
define view s3:Restaurant as  
select * from Restaurant
```

Connecting the source schemas to the global schema: Local-as-View



GAV and LAV have different expressive power

- **Some GAV scenarios cannot be expressed in LAV**
- Example:
create view **ParisPackage** as
select ph.name as hotelName, ph.street as hotelAddress,
r.name as restaurantName, r.rating as restaurantRating
from s1:ParisHotel ph, s3:Restaurants r
where **r.city='Paris' and r.street=ph.street**
- The view only contains (hotel, restaurant) pairs that are on the same street
- It is not possible to express this with LAV mappings
 - LAV describes each source *individually* w.r.t. the global schema
 - Not in correlation with data available in *other sources*

GAV and LAV have different expressive power

- There exist **LAV scenarios that cannot be expressed in GAV**
- Example:
 - Not possible to express in GAV that s1:ParisHotels only has data about Paris, while s2:LyonHotel only has data from Lyon
 - A query about hotels in Paris will also be sent to s2, although it will bring no results
 - LAV query processing avoids this (see next)

GAV and LAV have different expressive power

- There exist **GAV** scenarios that cannot be expressed in LAV
- Example:

```
create view ParisPackage as
select ph.name as hotelName, ph.street as hotelAddress, r.name as
restaurantName, r.rating as restaurantRating
from s1:ParisHotel ph, s3:Restaurants r
where r.city='Paris' and r.street=ph.street
```

- The closest we can do is define s1.ParisHotel and s3.Restaurants *each* as a projection over ParisPackage
- But this changes the semantics of ParisPackage:
 - It does not express that *only Paris restaurants* are in ParisPackage
 - Not possible to express that only (hotel, restaurants) *on the same street* are available through the integration system
 - ParisPackage becomes the cartesian product of ParisHotel with all restaurants...

Query processing in Local-as-View (LAV)

```
define view s1:ParisHotels as  
select street, name, price as roomPrice  
from Hotel where city='Paris'  
define view s2:LyonHotel as  
select street, name, descr as roomDesc, price as roomPrice  
from Hotel where city='Lyon'  
define view s3:Restaurant as  
select * from Restaurant
```

Query:

```
select h.street, h.price, r.rating  
from Hotel h, Restaurant r  
where r.city=h.city and h.street=r.street
```

Query processing in Local-as-View (LAV)

```
define view s1:ParisHotels as  
select street, name, price as roomPrice  
from Hotel where city='Paris'
```

```
define view s2:LyonHotel as  
select street, name, descr as roomDesc, price as  
roomPrice from Hotel where city='Lyon'
```

```
define view s3:Restaurant as  
select * from Restaurant
```

Step 1: identify
potentially useful
views

Query:

```
select h.street, h.price, r.rating from Hotel h, Restaurant r  
where r.city=h.city and h.street=r.street
```

Query processing in Local-as-View (LAV)

Query:

select h.street, h.price, r.rating from Hotel h, Restaurant r
where r.city=h.city and h.street=r.street

Step 2: generate **view combinations** that may be used to answer the query (one view for each table in the query):

s1:ParisHotels and s3:Restaurant

s2:LyonHotels and s3:Restaurant

Step 3: for each view combination and each view, check:

If the view returns the attributes we need:

Those returned by the query, *and*

Those on which possible query joins are based

If the view selections (if any) are compatible with those of the query

If one condition is not met, discard the view combination.

define view s1:ParisHotels as
select street, name, price as roomPrice
from Hotel where city='Paris'

The query needs:

- street, price, rating (returned): the view provides them
- city and street for the join: street is provided, city is not (but it is a constant, thus known)

The view has a selection on the city which the query did not have → The view provides part of the query result, but does not contradict the query.

The view s1:ParisHotels is OK.

define view s3:Restaurant as select * from Restaurant

The view s3:Restaurants is OK.

The view combination s1:ParisHotels, s3:Restaurants is OK provided that Restaurant.city is set to Paris.

Query processing in Local-as-View (LAV)

Query:

```
select h.street, h.price, r.rating from Hotel h, Restaurant r
where r.city=h.city and h.street=r.street
```

Step 2: generate **view combinations** that may be used to answer the query (one view for each table in the query):

s1:ParisHotels and s3:Restaurant

s2:LyonHotels and s3:Restaurant

Step 3: for each view combination and each view, check:

[...]

If one condition is not met, discard the view combination.

Step 4: for each view combination, add the necessary joins among the views, possibly selections and projections → rewriting

Query rewriting using s1:ParisHotels and s3:Restaurant:

```
select h.street, h.price, r.rating
from s1:ParisHotels h and s3:Restaurant r
where r.city='Paris' and h.street=r.street
```

This is a partial rewriting, and so is:

Query rewriting using s2:LyonHotel and s3:Restaurant:

```
select h.street, h.price, r.rating
from s2:LyonHotels h and s3:Restaurant r
where r.city='Lyon' and h.street=r.street
```

Query processing in Local-as-View (LAV)

Query:

select h.street, h.price, r.rating from Hotel h, Restaurant r
where r.city=h.city and h.street=r.street

Step 2: generate **view combinations** that may be used to answer the query (one view for each table in the query):

s1:ParisHotels and s3:Restaurant

s2:LyonHotels and s3:Restaurant

Step 3: for each view combination and each view, check:

[...]

If one condition is not met, discard the view combination.

Step 4: for each view combination, add the necessary joins among the views, possibly selections and projections → rewriting

Step 5: return the union of the rewritings thus obtained

Full query rewriting:

select h.street, h.price, r.rating
from **s1:ParisHotels** h and **s3:Restaurant** r
where r.city='Paris' and h.street=r.street
union all
select h.street, h.price, r.rating
from **s2:LyonHotel** h and **s3:Restaurant** r
where r.city='Lyon' and h.street=r.street

Query processing in Local-as-View (LAV)

```
define view s1:ParisHotels as... from Hotel where city='Paris'  
define view s2:LyonHotel as... from Hotel where city='Lyon'  
define view s3:Restaurant as select * from Restaurant
```

Query:

```
select h.street, h.price, r.rating  
from Hotel h, Restaurant r  
where r.city=h.city and h.street=r.street
```

Rewriting of the query using the views:

```
select h1.street, h1.price, r3.rating  
from s1:ParisHotels h1, s3:Restaurant r3  
where h1.city=r3.city and h1.street=r3.street
```

union all

```
select h2.street, h2.price, r3.rating  
from s2:LyonHotels h2, s3:Restaurant r3  
where h2.city=r3.city and h2.street=r3.street
```

Concluding remarks on Local-as-View (LAV)

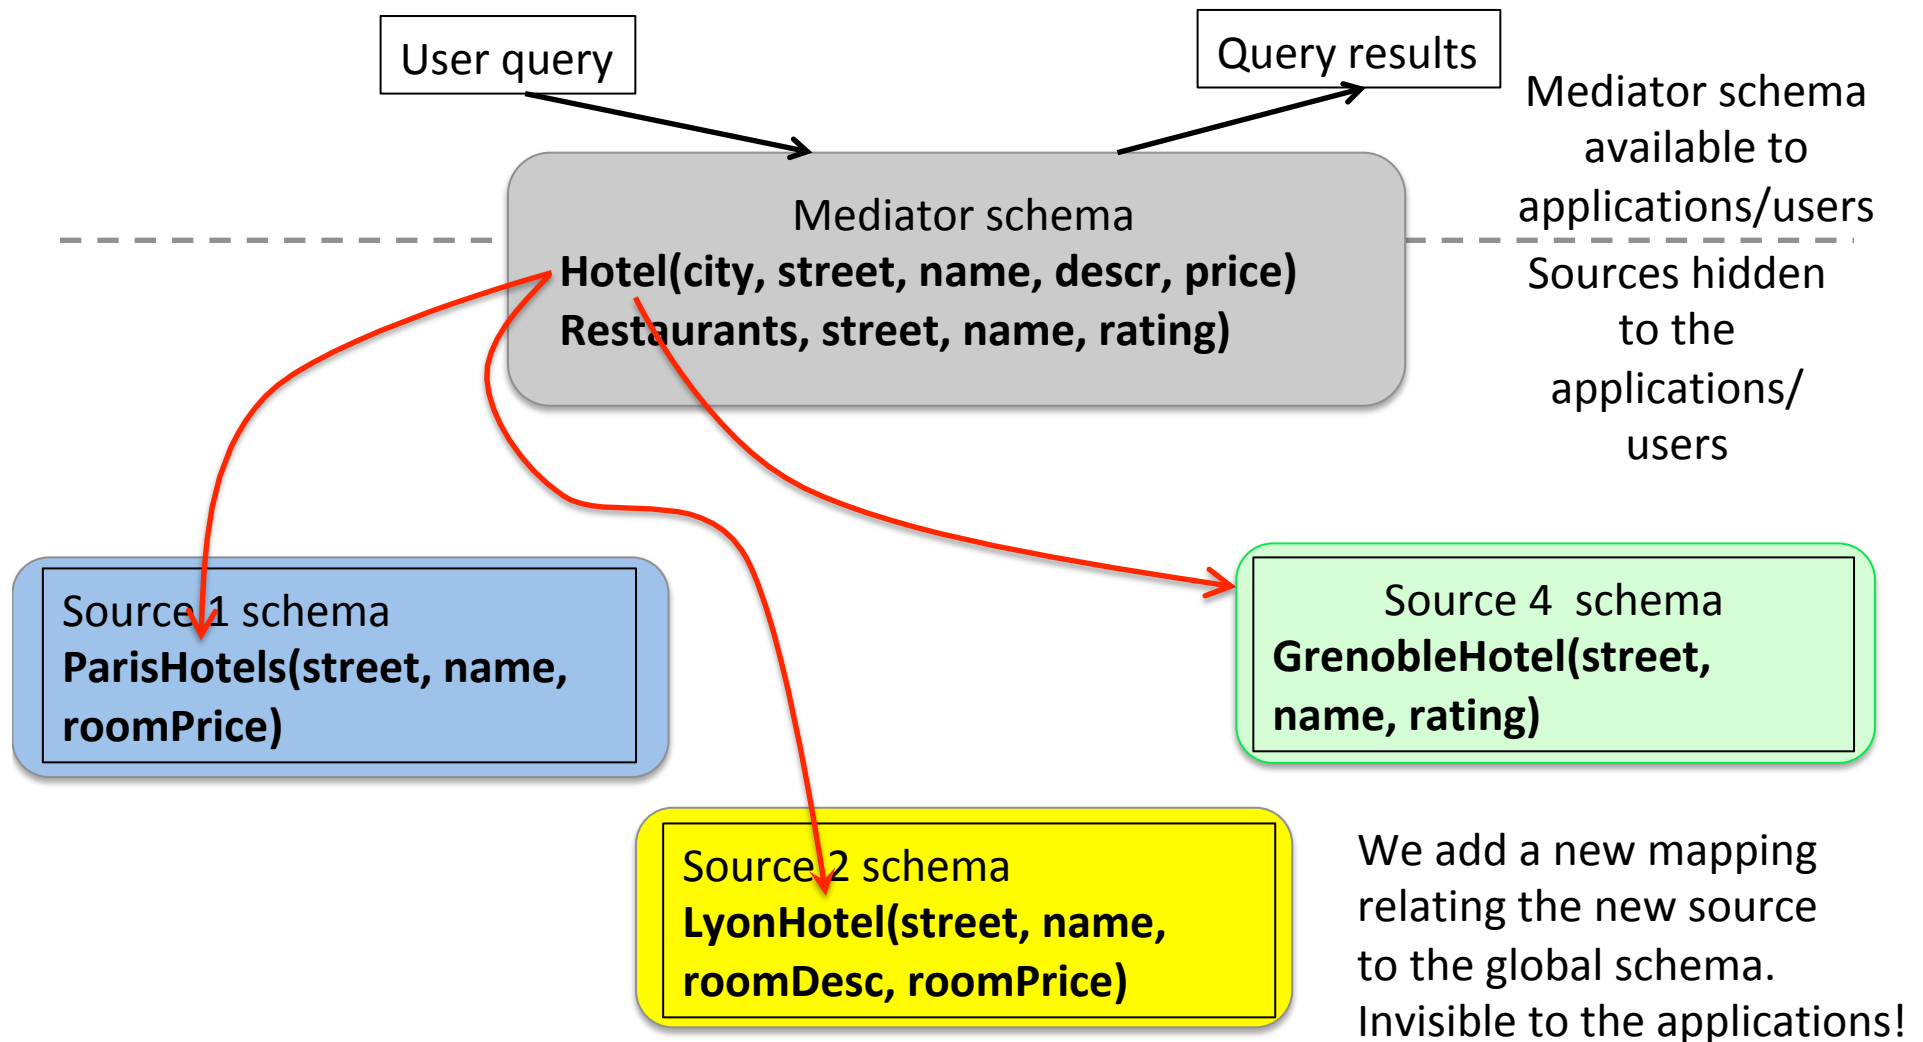
Query processing

- The problem of finding all rewritings given the source and global schemas and the view definitions = **view-based query rewriting**, NP-hard in the size of the (schema+view definitions).
 - These are often much smaller than the data

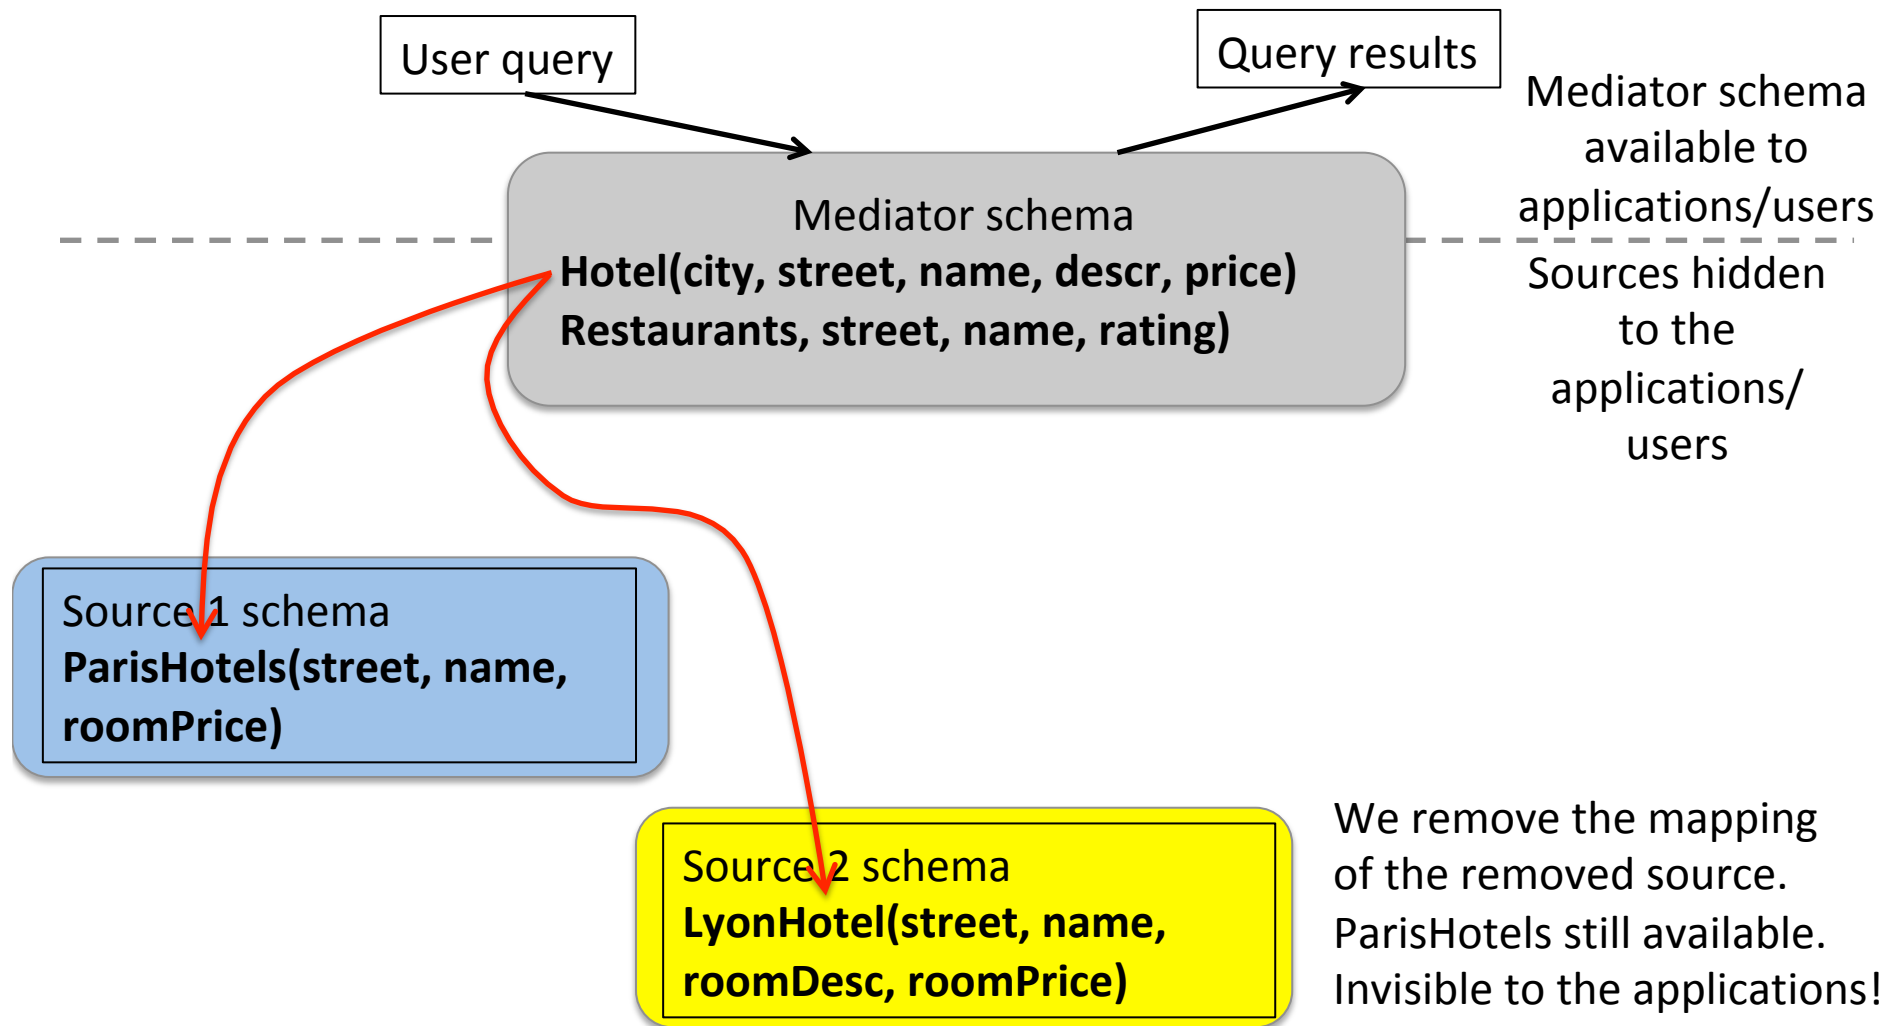
The schema definition is **more robust**:

- One can independently add/remove sources from the system without the global schema being affected at all (see next)
- Thus, no application needs to be aware of the changes in the schema

Local-as-View: adding a new source



Local-as-View: Removing a source



Connecting the source schemas to the global schema: Global-Local-as-View (GLAV)

Generalizes both GAV and LAV

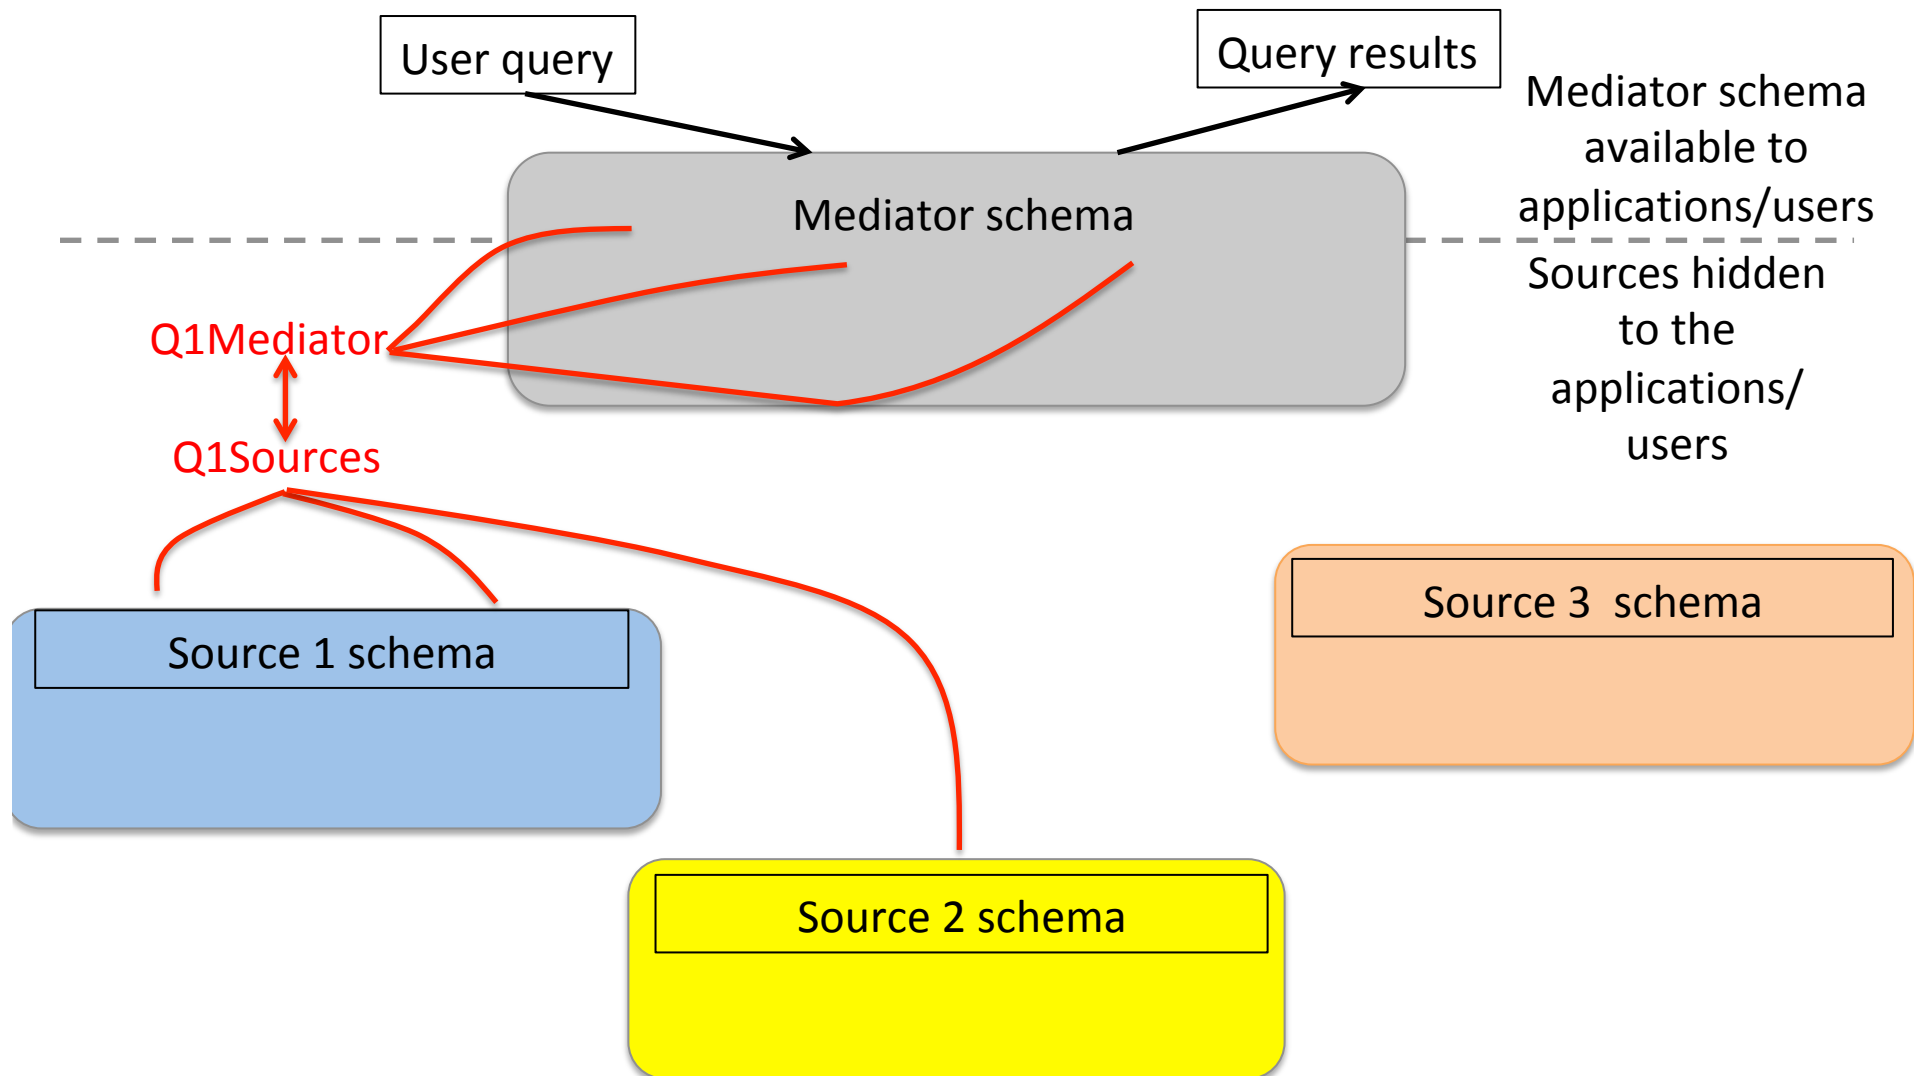
1 mapping = 1 pair (query over 1 or several sources schemas,
query over the mediator schema)

$$\begin{aligned} Q1Mediator(m:r1, m:r2, m:r3, \dots) &\leftrightarrow Q1Sources(s1:t1, s2:t1, \dots) \\ Q2Mediator(m:r1, m:r2, m:r3, \dots) &\leftrightarrow Q2Sources(s1:t1, s2:t1, \dots) \\ Q2Mediator(m:r1, m:r2, m:r3, \dots) &\leftrightarrow Q3Sources(s1:t1, s2:t1, \dots) \end{aligned}$$

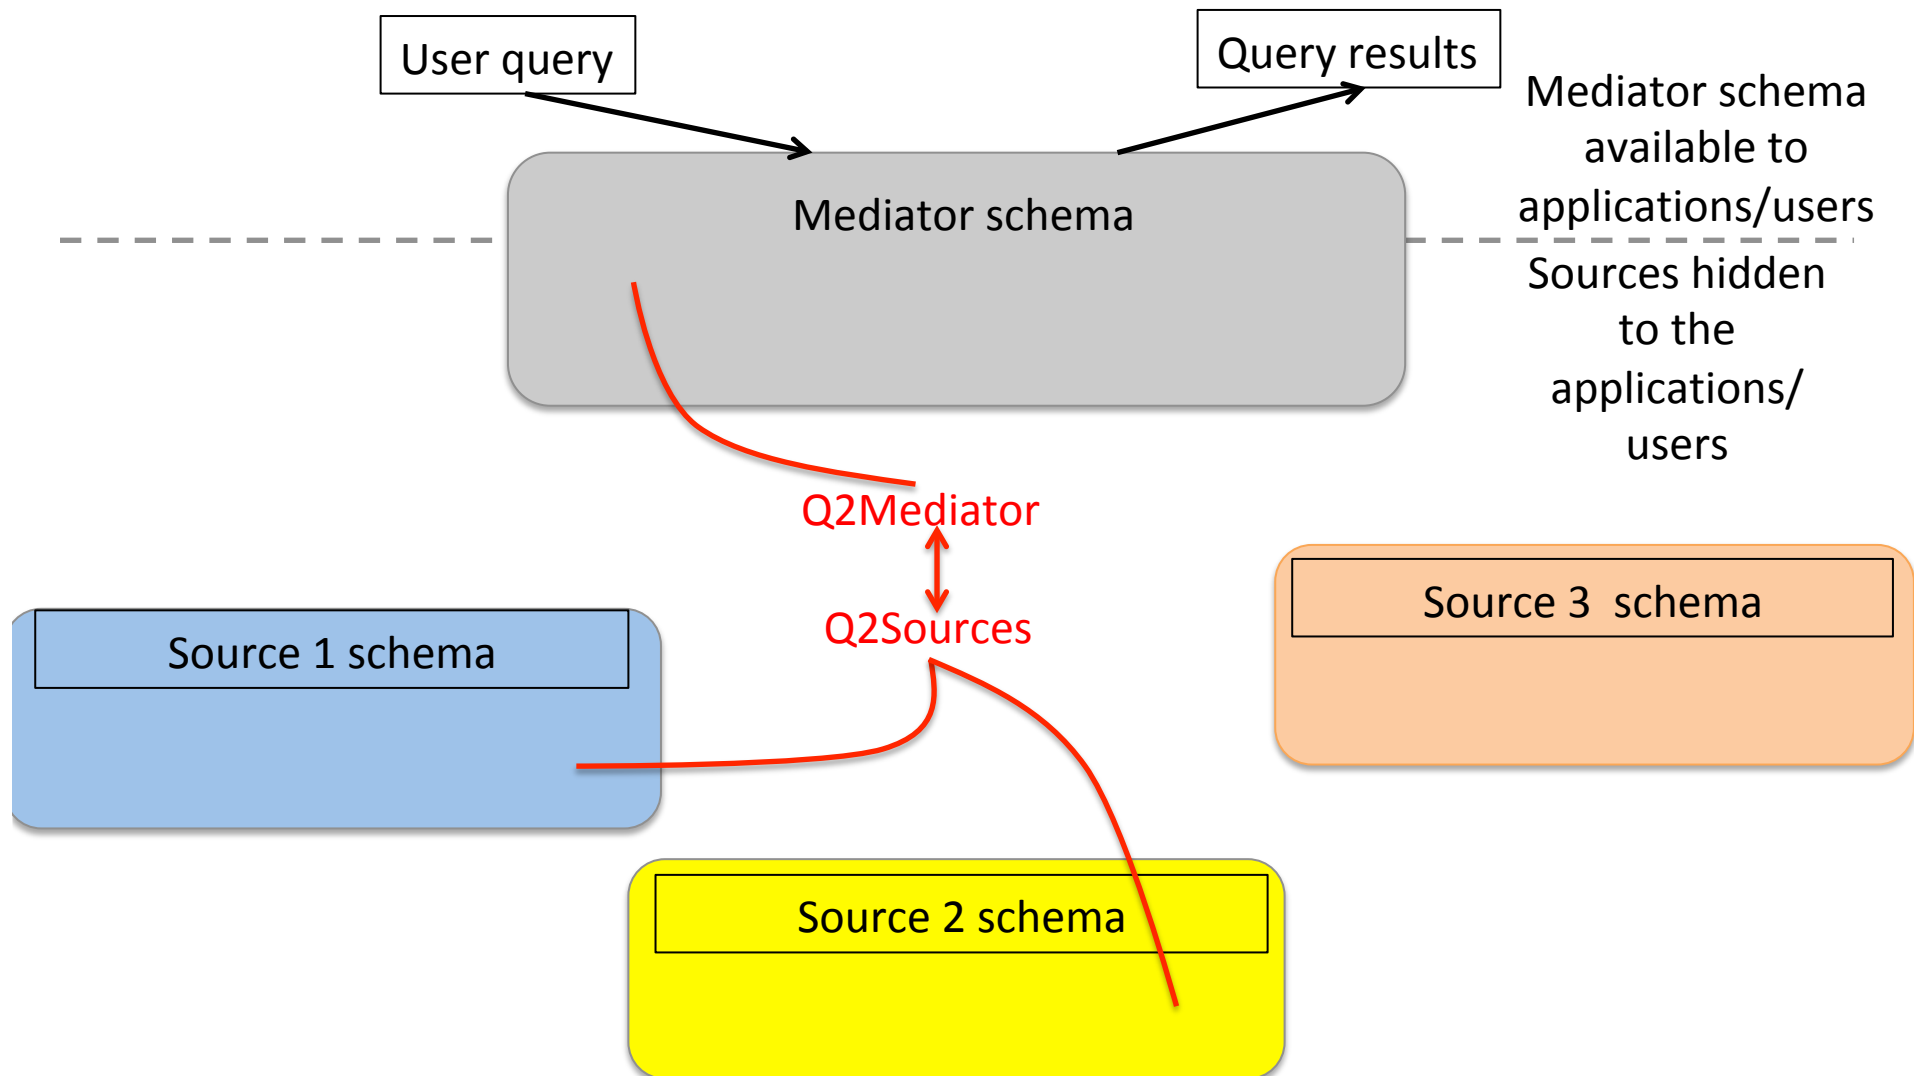
Semantics: there is a tuple in $Q_iMediator(\dots)$ for each result of $Q_iSources(\dots)$

- A GAV mapping is a particular case of GLAV mapping where $QMediator$ is exactly one mediator relation
- A LAV mapping is a particular case of GLAV mapping where $QSources$ is exactly one source relation

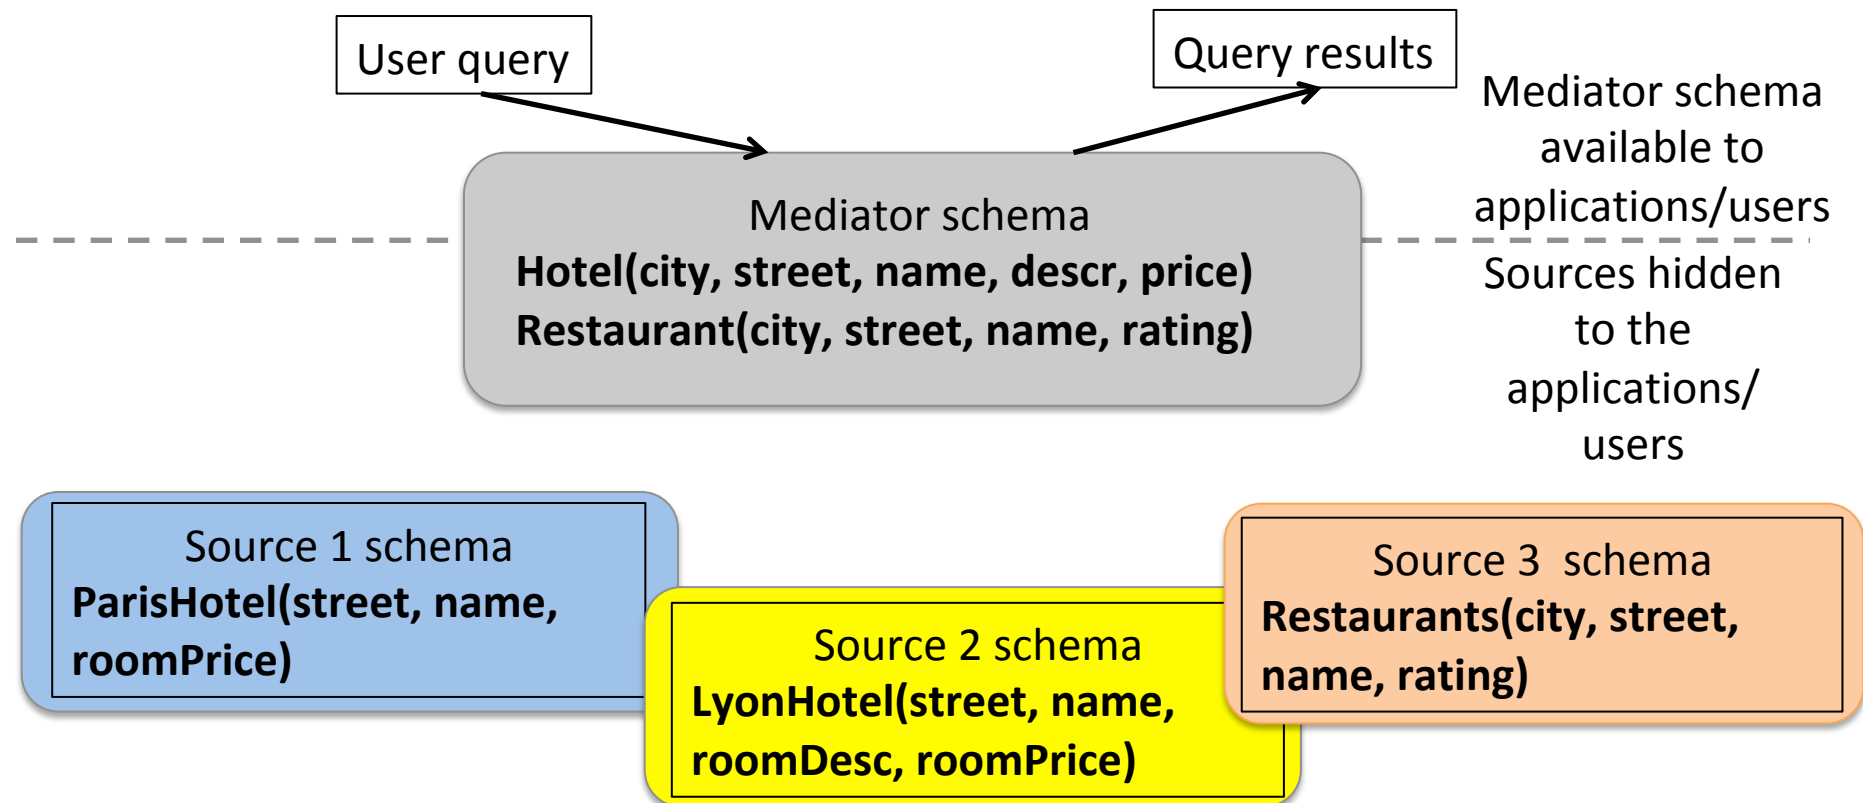
Connecting the source schemas to the global schema: Global-Local-as-View (GLAV)



Connecting the source schemas to the global schema: Global-Local-as-View (GLAV)



Global-Local-as-View: example

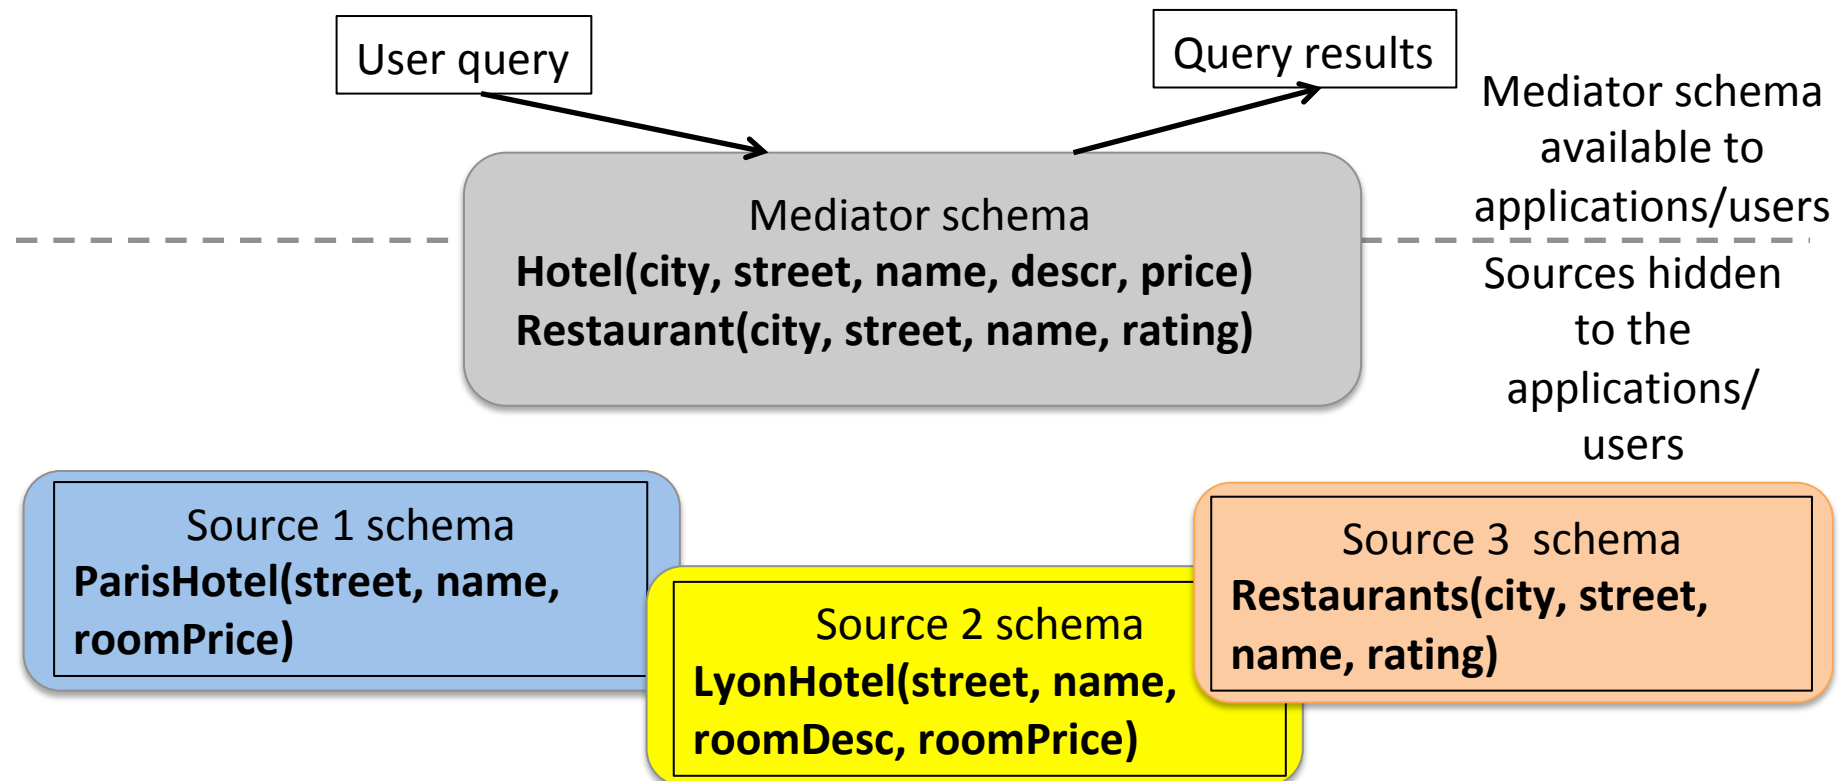


Previous LAV mapping of Source 1:

Q1Mediator: select street, name, price as roomPrice from Hotel where city='Paris'

Q1Sources: select * from ParisHotel

Global-Local-as-View: example



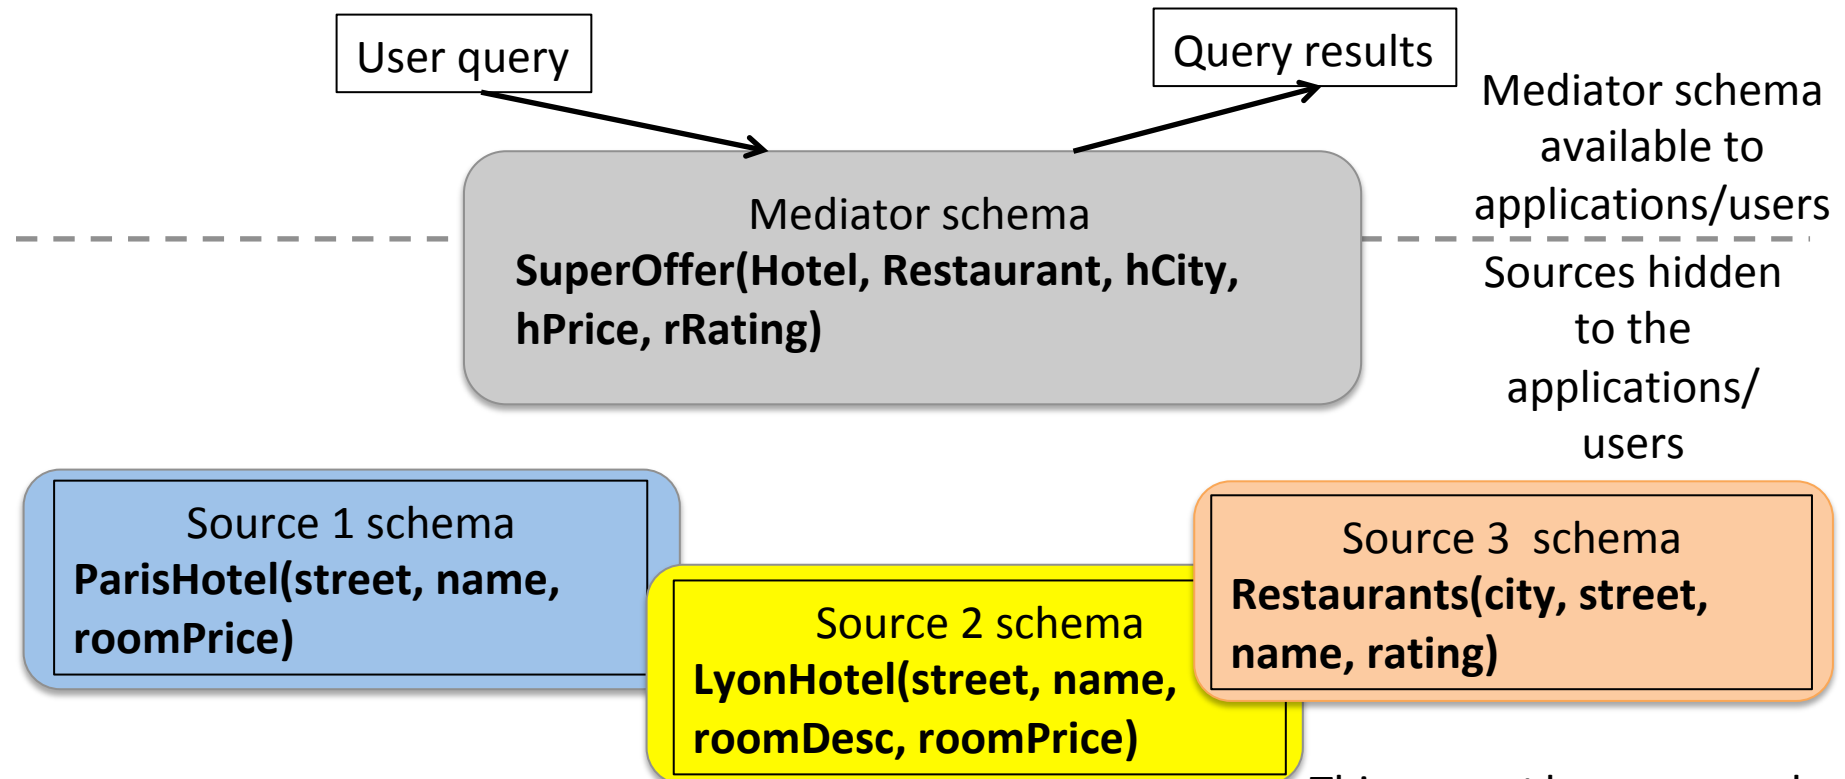
Previous GAV mapping of Hotel:

Q2Mediator: select * from Hotel

Q2Sources: select 'Paris' as city, street, name, null as descr, roomPrice as price from ParisHotel
union

select 'Lyon' as city, street, name, roomDesc as descr, roomPrice as price from LyonHotel

Global-Local-as-View: example



New GLAV mapping:

Q3Mediator: select * from SuperOffer where hCity='Lyon'

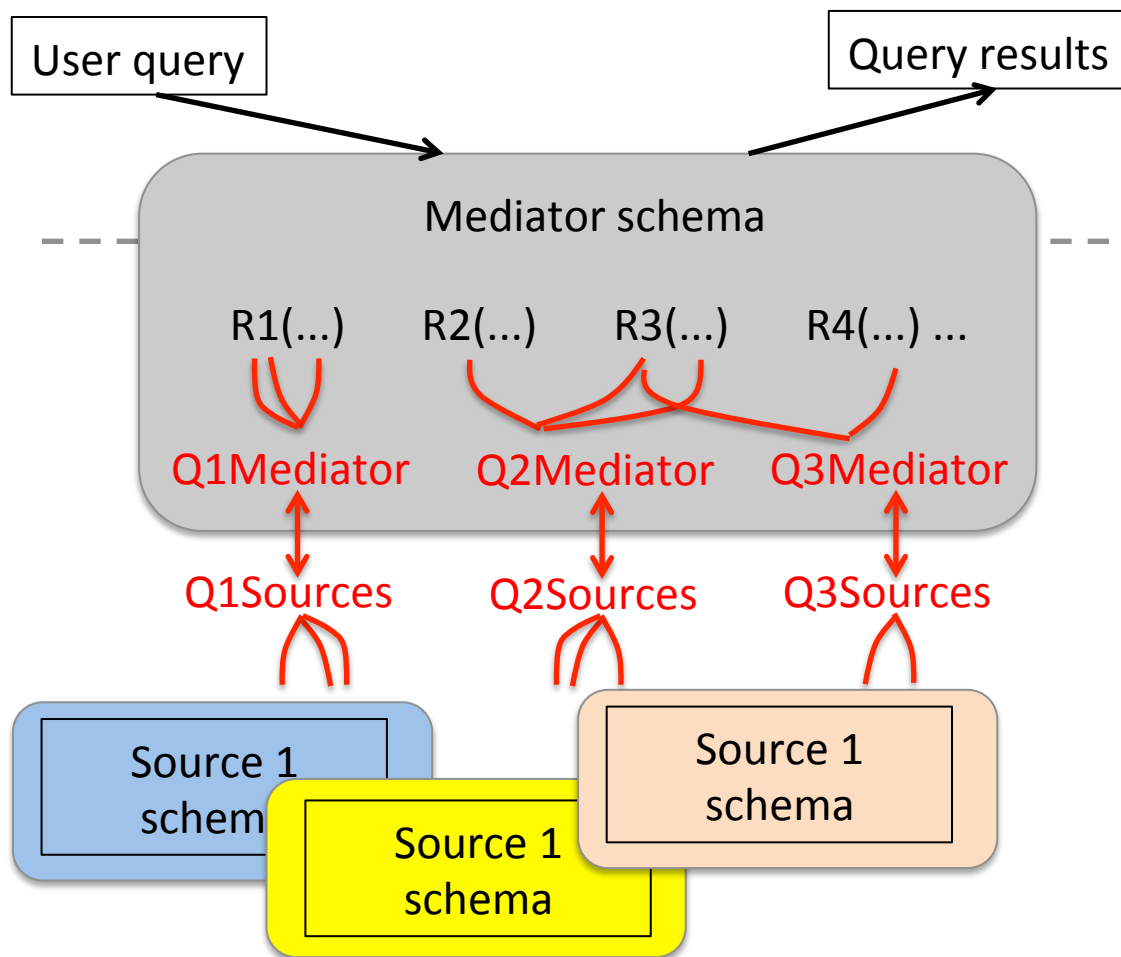
Q3Sources: select lh.name, r.name, h.roomPrice * 0.5 as hPrice, r.rating as rRating
from LyonHotel lh, Restaurants r
where r.city='Lyon' and name='Lion d'Or' and r.street=lh.street

This *cannot* be expressed
either in LAV or GAV.

This mapping says: "each result of Q3Sources leads to a SuperOffer in Lyon".

B Other mappings could define more SuperOffers in Lyon, or in other cities, or with rRating=3...

Query Processing in GLAV



User queries asked on the mediator schema.

Q1Mediator, Q2Mediator, ... are queries over this schema

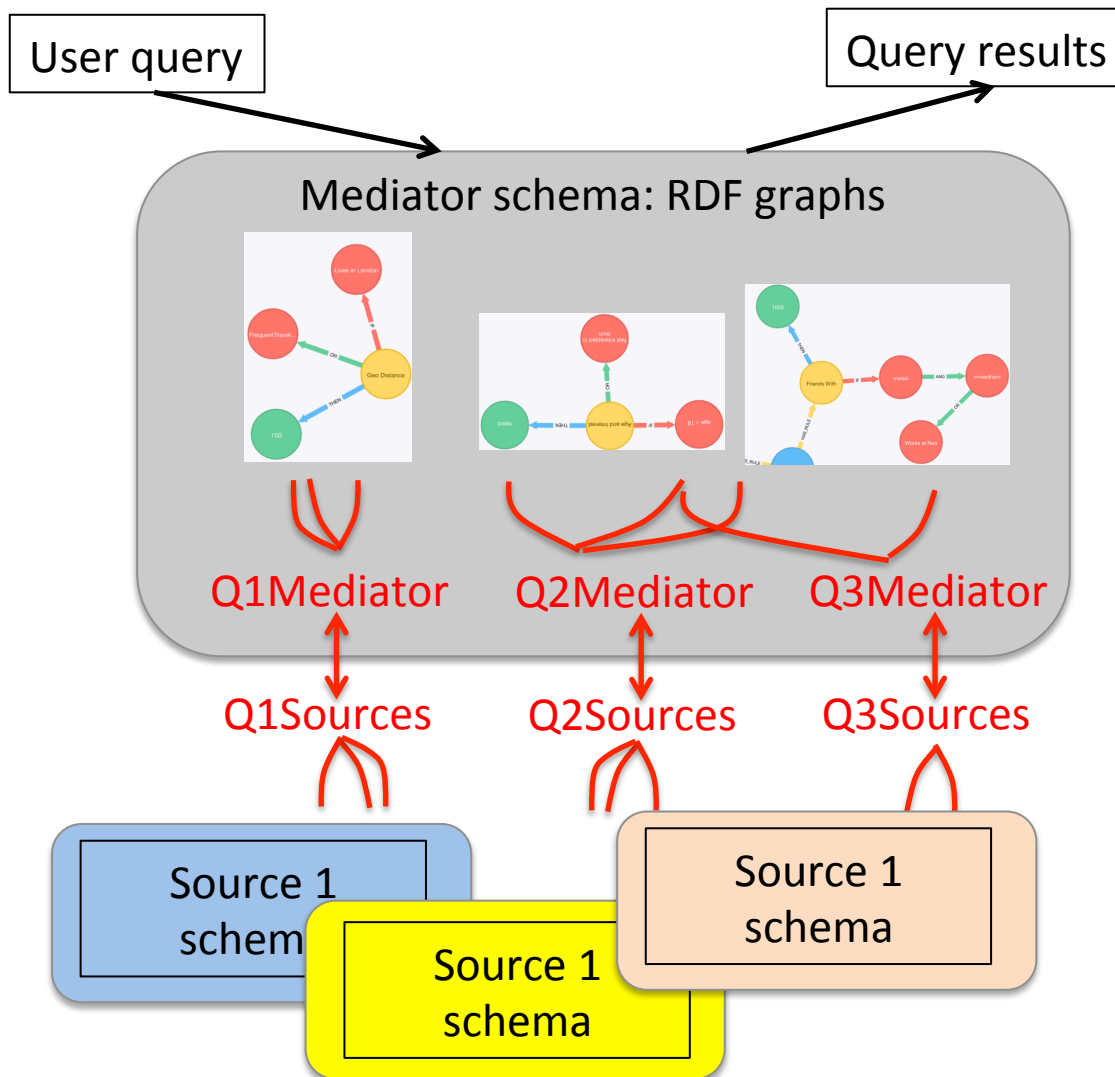
1. Apply **LAV**-style rewriting considering each QiMediator as a view over the mediator schema.
 - This leads to rewritings of Q over QiMediator relations (Q1Mediator, Q2Mediator, ...)
2. For each such rewriting, in **GAV** style, replace the symbol QiMediator by the query QiSources.
 - This leads to a rewriting of the query over the sources.

Examples: find all super offers in Paris? in Lyon?

Concluding remarks on GLAV

- The most flexible approach
 - Can express LAV, GAV, and more
- If a source changes or sources are added, as long as Q1Sources can be rewritten, applications will not be impacted
 - Only the "invisible" part of the system needs maintenance
- Query rewriting remains expensive because it includes view-based query rewriting (NP-hard) as well as query unfolding (very simple)

Modern mediators: GLAV with RDF global schema



Idea 1: RDF global schema

Advantages:

We do not have to fix a set of relations in advance

We can use ontologies to add semantics to the global schema

Idea 2: write GLAV mappings, e.g.:

1. **Q1Sources**: an SQL query returning (x, y, z) tuples

Q1Mediator: (x, 'friend', y), (y, 'worksfor' z)

Q1Mediator "creates RDF out of relational data"

2. **Q2Sources**: an XPath query returning (z) nodes

Q2Mediator:
(z, 'type', Company)

If common z value, the graphs built by Q1,2Mediator **connect**!

Concluding remarks on mediators

- **Data integration:** treat several data sources as a single one
 - Old problem that is not going away (quite the contrary)!

- **Needs:**


1. *Understand* the sources and *how they relate to the global schema* we want

2. Then, either:

1. *Extract* the data from the sources, *transform* it into the global schema, and *load* it into a **data warehouse (ETL)**, or

2. Devise a **mediator** which interacts with the sources and provides the illusion of a single database.

We have seen GAV, LAV and GLAV mediation.



DATA SPACES, DATA LAKES

Data spaces

- "Data spaces" (Franklin, Halevy, Maier, 2005):
 - Many heterogeneous data sources...
 - On a single or on multiple machines
 - But, unlike data integration systems, the sources
 - May not be **structured**: text, email, Web pages, directories...
 - Therefore, different data models, or unstructured (no data model) = text
 - May not reside in **databases**
 - Therefore, no source query language

Data spaces

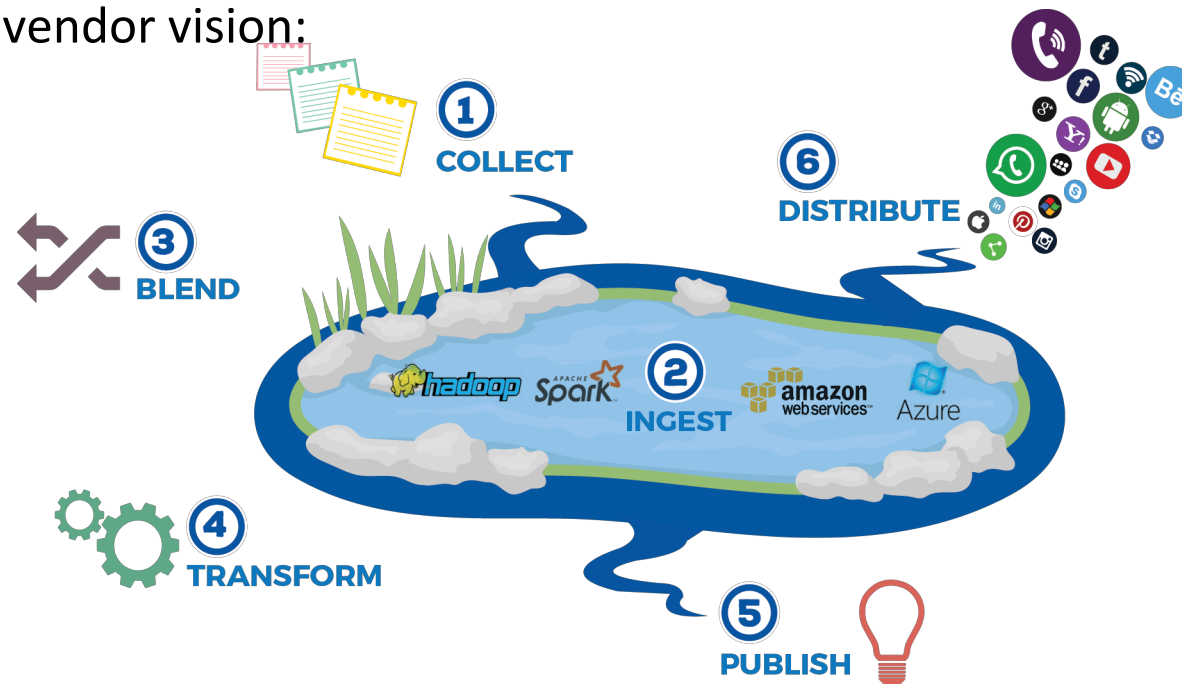
- How to query the data space?
- Keyword query: kw1, kw2, ..., kwn
 - E.g., "rent flat Chamonix"
- Answers:
 - From a text file: **minimal text fragments** that contain all kwds
 - From a database:
 - **One tuple** if it contains all the kwds, or
 - **A few tuples** if they join and they contain all the kwds, or
 - **A minimal JSON tree** that contains all the kwds etc.
 - *Score* to decide which answers to return first
- **No schema, no integration effort:** too many sources, too heterogeneous
- Keyword query paradigm

Data lakes

- Popular term, started around 2010 (cca)
- Mostly in **companies**
- Many data sources: Tens, hundreds, thousands
 - Most of the time relational. Also: text, JSON
 - Developed more or less independently of each other, with no knowledge of each other
 - Different schemas; different names for same things; slightly different semantics (e.g., "customer" vs. "customer who bought something in the last year")
 - Some relationships probably exist between the schemas of the different databases
 - ... but finding and expressing them has become beyond current human capacity

Data lake: usage

- Positive vendor vision:



- The hard part is BLEND because this requires understanding data which...
 - Has been designed 10 years ago by someone who has since left the company...
 - Was meant for (or was gathered by) an application the company no longer uses..
 - Lacks documentation (or the documentation obsolete)...
 - Overlaps partially with a few other sources and (it is feared) with many others...
- No point in learning from data we don't understand!

Data lakes: problems and products

- Problems:
 - Automatically **summarizing** a data source: *data profiling*
 - **Identifying relationships** between different data sources: *data matching, data profiling, data cleaning*
 - So that the data lake is not a "data swamp"
 - Build an understanding/relationships between the data sources over time
 - Query processing over data sources whose relationships are well understood follow the mediator or the warehouse (ETL) path

Data lake products:

- <https://www.ibm.com/analytics/data-management/data-lake>
- <https://blogs.oracle.com/emeapartnerbiepm/oracle-analytics-cloud-data-lake-edition-available>