

# JSON stores

**Ioana Manolescu**

Inria and Ecole polytechnique

# Motivation


- JSON (JavaScript Object Notation) allows to describe nested, potentially heterogeneous data
  - Very flexible
  - Thus, a good idea for NoSQL!
- Less verbose than XML

# Sample JSON document: Twitter


```
{ "results":[
  { "text": "@twitterapi http://tinyurl.com/ctrefg",
    "to_user_id": 396524,
    "to_user": "TwitterAPI",
    "from_user": "jkoum",
    "metadata": { "result_type": "popular", "recent_retweets": 109 },
    "id": 1478555574,
    "from_user_id": 1833773,
    "iso_language_code": "nl",
    "source": "twitter< /a>",
    "profile_image_url": http://s3.amazonaws.com/twitter/a155\_b\_normal.jpg,
    "created_at": "Wed, 08 Apr 2009 19:22:10 +0000"},
    ... truncated ... ],
  "refresh_url": "?since_id=1480307926&q=%40twitterapi",
  "results_per_page": 15,
  "next_page": "?page=2&max_id=1480307926&q=%40twitterapi",
  "completed_in": 0.031704,
  "page": 1,
  "query": "%40twitterapi"
}
```

# JSON document structure

- Object: collection of (name, value) pairs



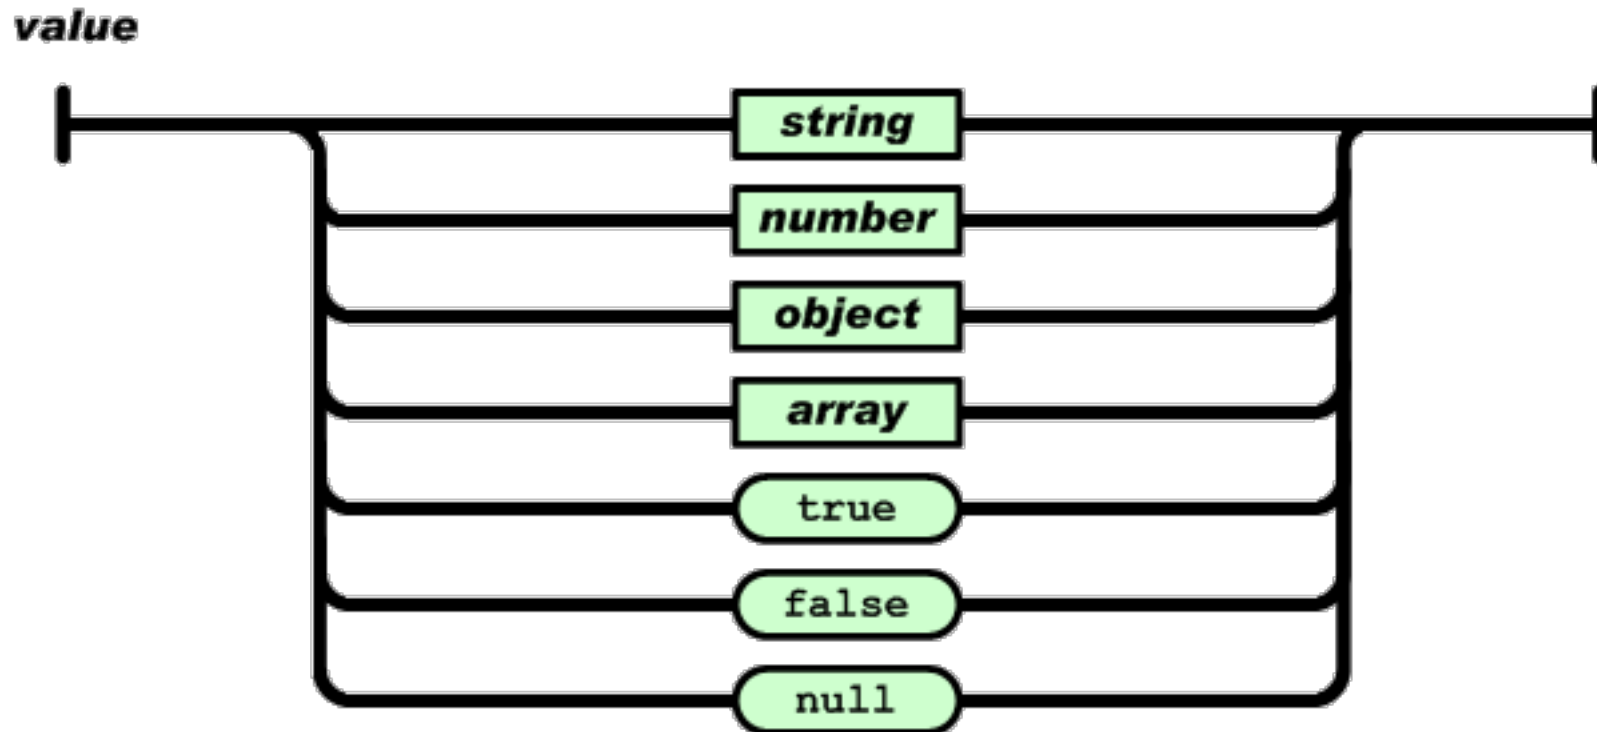
- Array: collection of values



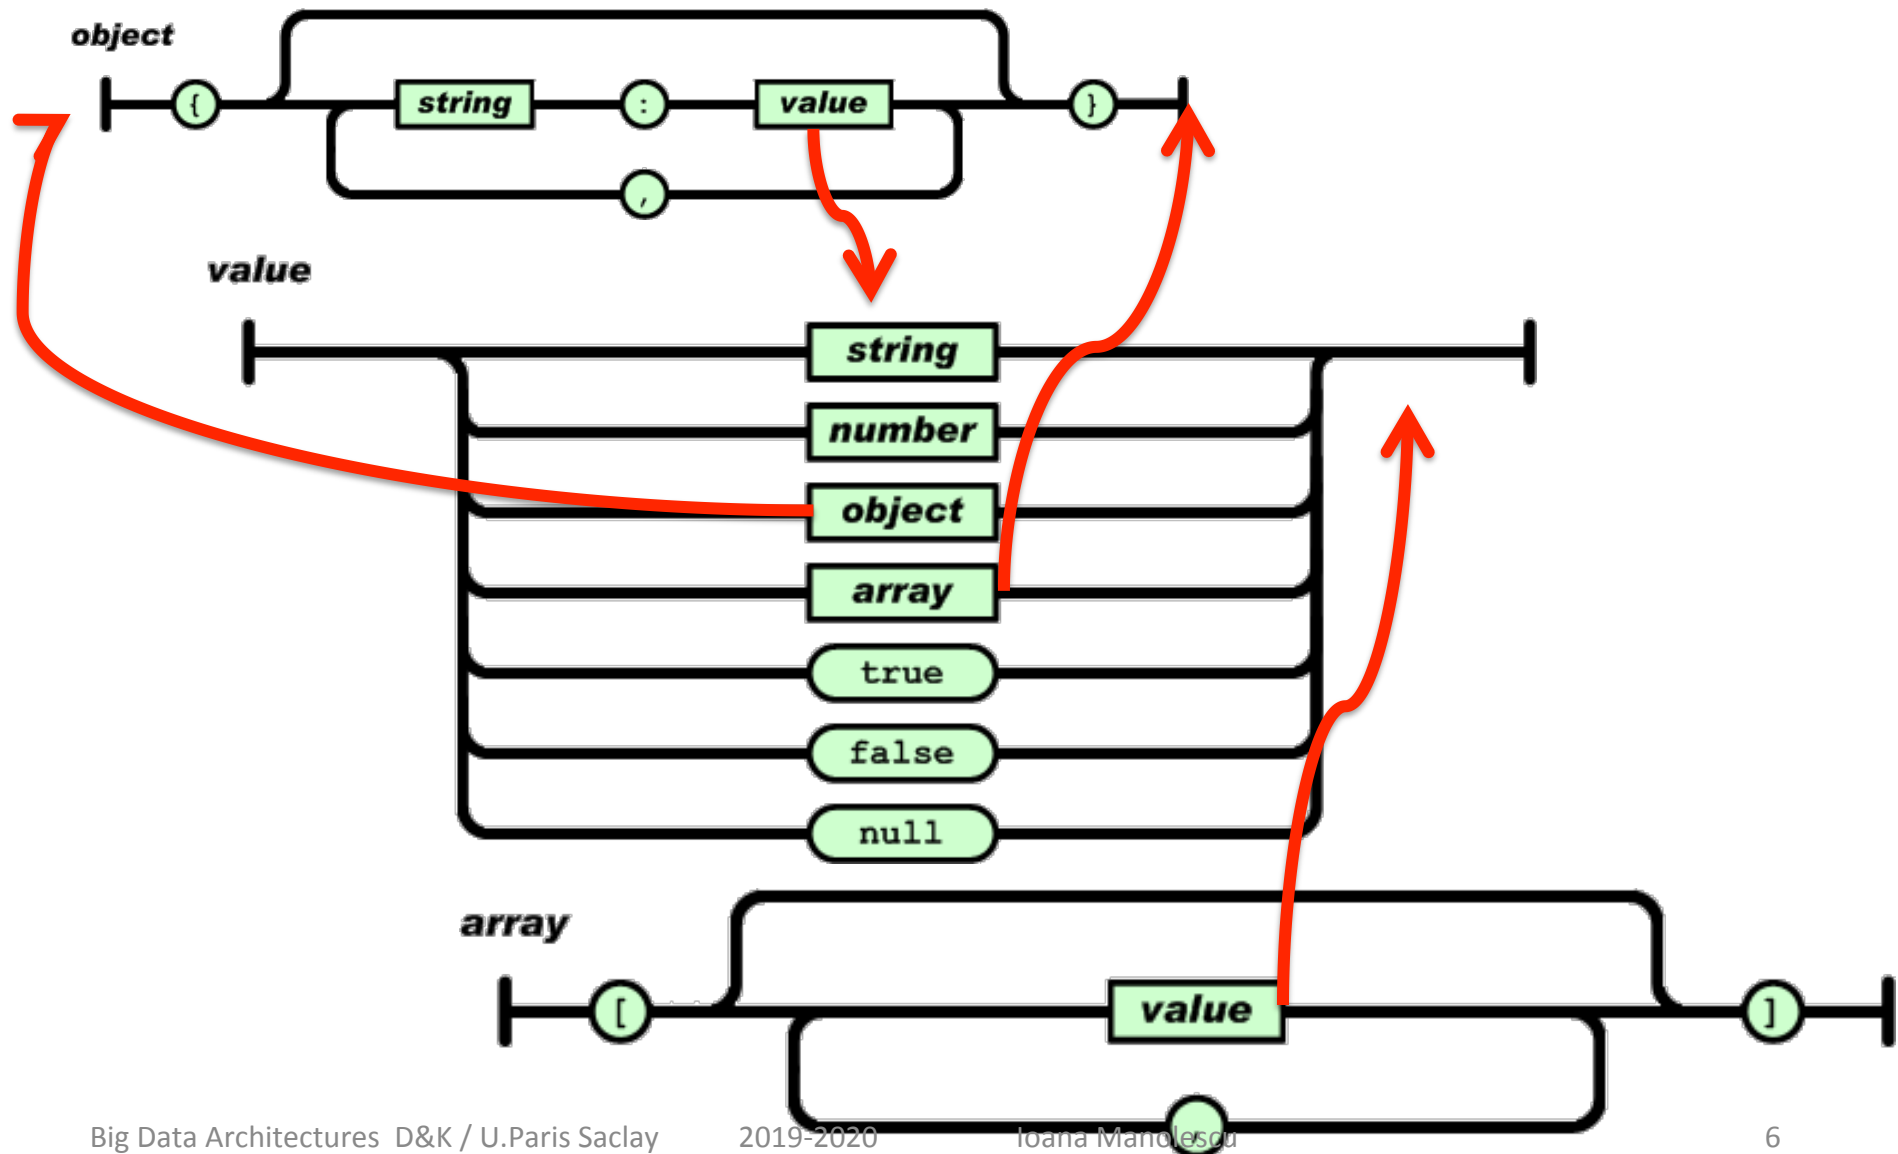
- Arrays and object structure are heterogeneous (no schema)

# JSON document structure


- Values (allow nesting):



# JSON document structure




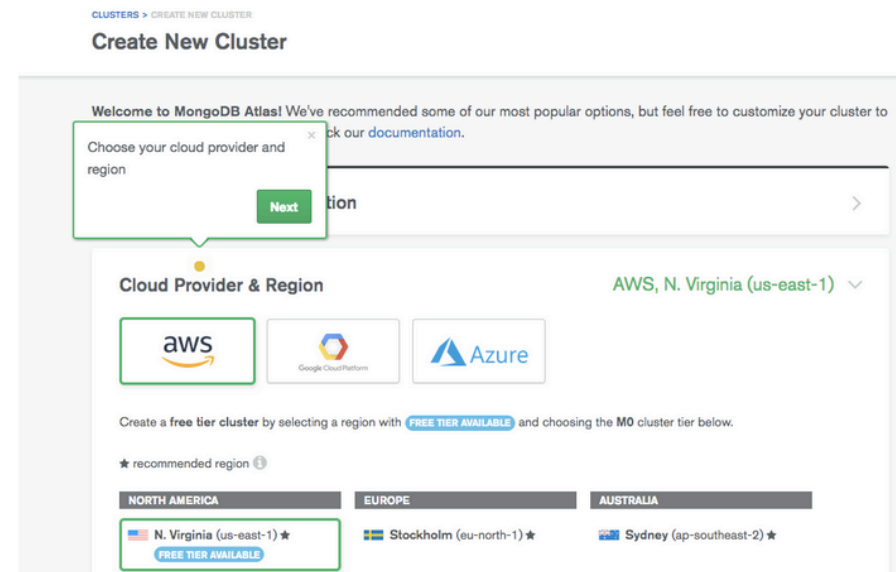
# MongoDB: a JSON document store



Computed based on: popularity in search engine results, queries, job offers,  
social networks, questions on StackOverflow...

# MongoDB tools

- **MongoDB server (mongod)**
  - Installed on top of an operating system (e.g., on a PC); assumes a file system and regular OS support
  - Accessed through (potentially concurrent) clients (mongo)
- **MongoDB Atlas (or cloud edition)**
  - Meant to deploy easily in the cloud (e.g. Amazon, Google, Azure...)
  - Assumes the respective cloud infrastructure services
- **Other MongoDB tools:** 
  - Data lake
  - Charts to generate reporting images
  - Compass to visualize the data
  - Spark connector








# MongoDB server

- The server data can be:
  - Replicated
    - Several identical copies of the same data
  - Partitioned (sharded)
    - Distributed across the machines of a cluster in order to take advantage of the storage and processing capacity
- Document processing
  - Selective access, or
  - Map-Reduce mode

# MongoDB storage organization

- **Documents** are stored in **collections** (which may have **indexes**)
  - Collections are part of **databases**
- ```
> mongo myExample           // Creates the database myExample   
> db.towns.insert({         // Creates the collection towns  
  name: "New York",         // and inserts a document into it  
  population: 22200000,  
  last_census: ISODate( "2009-07-31"),  
  famous_for: [ "statue of liberty", "food" ],  
  mayor : {  
    name : "Michael Bloomberg",  
    party : "I"  
  }  
}
```


# MongoDB object IDs

```
> db.towns.find()
{
   "_id" : ObjectId("4d0ad975bb30773266f39fe3"),
  "name" : "New York",
  "population": 22200000,
  "last_census": "Fri Jul 31 2009 00:00:00 GMT-0700
(PDT)",
  "famous_for" : [ "statue of liberty", "food" ],
  "mayor" : { "name" : "Michael Bloomberg", "party" : "I" }
}
```

\_id is implicitly added by the system  
Each object ID is different

# MongoDB object IDs

12 bytes:



Timestamp; client machine ID; process ID;  
incremented counter



IDs are unique across machines and databases

# MongoDB: information about data

It is possible to ask questions about objects, functions etc.

```
> typeof db
```

```
object
```

```
> typeof db.towns
```

```
object
```

```
> typeof db.towns.insert
```

```
function
```

# MongoDB: Operations

- Can find meta-informations about data (e.g., type)
- Allows CRUD operations:
  - Insert / **Create**
  - Query / **Read**
  - **Update**
  - **Delete**


# MongoDB: getting information about data

Calling a function with no parentheses shows the function code

**db.towns.insert**

```
function (obj, _allow_dot) {  
  if (!obj) { throw "no object passed to insert!"; }  
  if (!_allow_dot) { this._validateForStorage(obj); }  
  if (typeof obj._id == "undefined") {  
    var tmp = obj;  
    obj = {_id:new ObjectId};  
    for (var key in tmp) { obj[key] = tmp[key];}  
  }  
  this._mongo.insert(this._fullName, obj);  
  this._lastID = obj._id;  
}
```

# Working with (JavaScript) functions

Typing this into the client shell registers the function: 

```
function insertCity(name, population, last_census,  
                    famous_for, mayor_info) {  
    db.towns.insert( {name:name, population:population,  
                     last_census: ISODate(last_census),  
                     famous_for:famous_for,  
                     mayor : mayor_info}  
    );  
}
```

*Object implicitly created*



# Working with functions

Calling the function previously defined:

```
insertCity("Punxsutawney", 6200,  
'2008-31-01',  
["phil the groundhog"], {name: "Jim Wehrle" })
```

*Array*

*Object*



```
insertCity("Portland", 582000, '2007-20-09',  
["beer", "food"],  
{ name : "Sam Adams", party : "D" } )
```


db.towns.find() returns three objects

# Searching a MongoDB collection

```
db.towns.find({ "_id" :  
ObjectId("4d0ada1fbb30773266f39fe4") })
```

returns the full object

```
db.towns.find({"_id" :  
ObjectId("4d0ada1fbb30773266f39fe4") },  
{ name : 1 })
```

 This is called a projection operator


```
{  
  "_id" : ObjectId("4d0ada1fbb30773266f39fe4"),  
  "name" : "Punxsutawney"  
}
```

The `_id` field is the only one always returned by default (unless excluded, see next)

# Using exclusion in MongoDB projections

To exclude an attribute from a result, **set it to 0** in the find parameter

```
db.towns.find({ _id :  
ObjectId("4d0ada1fbb30773266f39fe4") }, { name : 0 })  
{  
  "_id" : ObjectId("4d0ada1fbb30773266f39fe4"),  
  "population" : 6200,  
  "last_census" : "Thu Jan 31 2008 00:00:00 GMT-0800  
(PST)",  
  "famous_for" : [ "phil the groundhog" ]  
}
```



A projection *cannot* contain *both* include and exclude specifications, except for the exclusion of the `_id` field.

In projections that *explicitly include* fields, only `_id` can be *excluded*.

# More search patterns

```
db.towns.find(  
  { name : /^P/, population : { $lt : 10000 } },  
  //selection  
  { name : 1, population : 1 })  
  // projection
```

→

```
{ "name" : "Punxsutawney", "population" : 6200 }
```

```
db.towns.find(  
  { last_census : { $lte : ISODate('2008-31-01') } },  
  { _id : 0, name: 1 })
```

→

```
{ "name" : "Punxsutawney" }  
{ "name" : "Portland" }
```

# Searching in nested structures

```
db.towns.find(  
  { famous_for : "food" },  
  { _id : 0, name : 1, famous_for : 1 })
```

→

```
{ "name" : "New York", "famous_for" : [ "statue of  
liberty", "food" ] }  
{ "name" : "Portland", "famous_for" : [ "beer", "food" ] }
```

```
db.towns.find(  
  { famous_for : { $all : ["food", "beer"] } },  
  { _id : 0, name:1, famous_for:1 })
```

\$all matches an array  
containing all the  
specified values

→

```
{ "name" : "Portland", "famous_for" : [ "beer",  
"food" ] }
```

# Searching in nested structures

Nodes which *must not* have a match of the search condition:

```
db.towns.find(  
  { famous_for : { $nin : ["food", "beer"] } },  
  { _id : 0, name : 1, famous_for : 1 })
```

→

```
{ "name" : "Punxsutawney",  
  "famous_for" : [ "phil the groundhog" ] }
```

Paths in conditions:

```
db.towns.find( { "mayor.party" : "I" }, { _id : 0, name : 1,  
mayor : 1 })
```


→

```
{ "name" : "New York",  
  "mayor" : { "name" : "Michael Bloomberg", "party" : "I" }  
}
```

# Searching in nested structures

Countries that export **bacon and tasty food**:

```
db.countries.find(
  { "exports.foods.name" : "bacon",
    "exports.foods.tasty" : true }, { _id : 0, name : 1 } )
```



Countries that export **tasty bacon**:

```
db.countries.find(
  { "exports.foods" : { $elemMatch : { name : "bacon",
                                       tasty : true } } },
  { _id : 0, name : 1 } )
```



Matched by:


```
{_id : "us", name : "United States",
 exports : { foods : [ { name : "bacon", tasty : true },
                       { name : "burgers" } ] }
}
```

# More search operators

- **\$regex**: matches PCRE-compliant regexes within `/... /`
- PCRE: Perl-Compatible Regular Expressions
- **\$ne, \$lt, \$lte, \$gt, \$gte**: arithmetics
- **\$exists, \$all, \$in, \$nin, \$or, \$nor, \$not**: logical operators
- **\$elemMatch**
- **\$size**: matches array of given size
- **\$type**: matches if field is of a given type



# MongoDB updates

```
db.towns.update(  
  { _id: ObjectID("4d0ada87bb30773266f39fe5") },   
  { $set : { "state" : "OR" } } );  
                                     // updates state
```

```
db.towns.update(  
  { _id: ObjectID("4d0ada87bb30773266f39fe5") },  
  { { "state" : "OR" } } );  
                                     // replaces the whole document!
```

```
db.towns.delete(  
  { _id: ObjectID("4d0ada87bb30773266f39fe5") } )  
                                     // deletes the document
```

# More operators used in updates

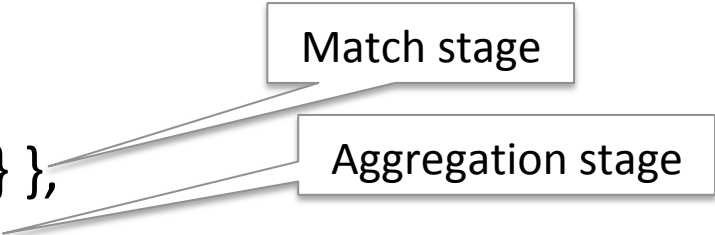
- **\$set, \$unset** (removes the field)
- **\$inc** (increments)
- **\$pop, \$push, \$pushall** for arrays
- **\$addToSet** like `push` but avoids duplicates
- **\$pull** removes a matching value from an array
- **\$pullAll** removes all matching values

# Other Operations: Aggregation

- **Aggregation** : can be applied as part of a pipeline consisting of several stages

```
db.collection.aggregate( [ { <stage> }, ... ] )
```

```
db.orders.aggregate([  
  { $match: { status: "A" } },  
  { $group: { _id: "$cust_id", total: { $sum: "$amount" } } }  
])
```



# Combining information from several documents (1/3)

MongoDB is mostly document (collection)-oriented

- Ideally, data which should be used together is stored in the same document(s)

To **combine** data from several documents, several options:

**1. Manual references:** store the ID of a document within another:

```
original_id = ObjectId()
```

```
db.places.insert({ "_id": original_id,  
                  "name": "Broadway Center",  
                  "url": "bc.example.net" })
```

```
db.people.insert({ "name": "Erin",  
                  "places_id": original_id,  
                  "url": "bc.example.net/Erin" })
```

# Combining information from several documents (2/3)

2. Use a **DBRef** (JSON object with "standard" attribute names). This is recognized as a pointer to another document:

```
{ "$ref" : <value>, "$id" : <value>, "$db" : <value> }
```

- "\$ref" points to the collection
- "\$id" points to the document
- "\$db" points to the database

```
{ "_id" : ObjectId("5126bbf64aed4daf9e2ab771"),    // ..  
  "creator" : { "$ref" : "creators",  
                "$id" : ObjectId("5126bc054aed4daf9e2ab772"),  
                "$db" : "users" } }
```

Recognized in some language drivers, not in all

E.g. Java, Python, Perl, PHP OK

E.g. C, C++, Scala: not supported (as of 4.2)

# Combining information from several documents (3/3)

Another way to combine data from several documents is a **lookup stage in a pipeline**.

A **\$lookup** stage implements a *left-outer join*, where one can filter documents from an input collection based on another collection

Syntax of a lookup stage:

```
{  
  $lookup: { from: <collection to join>,  
             localField: <field from the input documents>,  
             foreignField: <field from the documents of  
                           the "from" collection>,  
             as: <output array field> } }  
}
```

Lookups performs an **equality match** between localField and foreignField



Lookups outputs the documents **from the local collection** that survived the join (left outer join semantics)

# Global processing with custom code

One can define a JS function and run it

```
db.towns.find(  
  ( function() { return this.population >  
60000; } );  
  // runs the function over all the towns
```

Fails if one town has no population! 

Contrast with XML/XQuery

"OK for extra, not OK for missing"

# Indexing MongoDB data

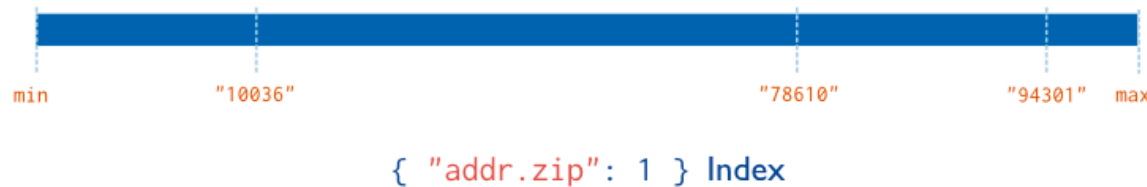

Indexes can be built on a collection calling `collection.createIndex({attr...})`

## 1. B-trees for exact and inequality search

- May be built on a single attribute (simple) or several attributes (compound)
- `Collection.createIndex({"name": 1})` ; *// for ascending order, otherwise use -1*
- B-tree index automatically built on `_id`

## 2. Multikey indexes allow indexing on an array attribute

- Built by default when one requires the indexing of an array attribute





# Indexing MongoDB data

## 3 **Geospatial indexes: 2d** (planar geometry based on x,y ), **2d sphere** (latitude, longitude)

- Operator: **\$near**, **\$nearsphere** (coordinates) returns the top k closest documents to the given coordinates
- Operators: **\$geoWithin**, **\$geoIntersects** (JSON rectangle)

## 4 **Text indexing** for full-text search

`db.reviews.createIndex( { comments: "text" } )`

`db.collection.createIndex( { "$**": "text" } )`

Inspect indexes: `db.system.indexes.find()`

# Searching with and w/o an index

## Without an index:

```
db.phones.find({display:
"+1800-5650001"})
.explain()
```

```
{ "cursor" : "BasicCursor",
  "nscanned" : 109999,
  "nscannedObjects" :
    109999,
  "n" : 1,
  "millis" : 52,
  "indexBounds" : { }
}
```

## With an index on display:

```
db.phones.find({ display:
"+1800-5650001" }).
explain()
```

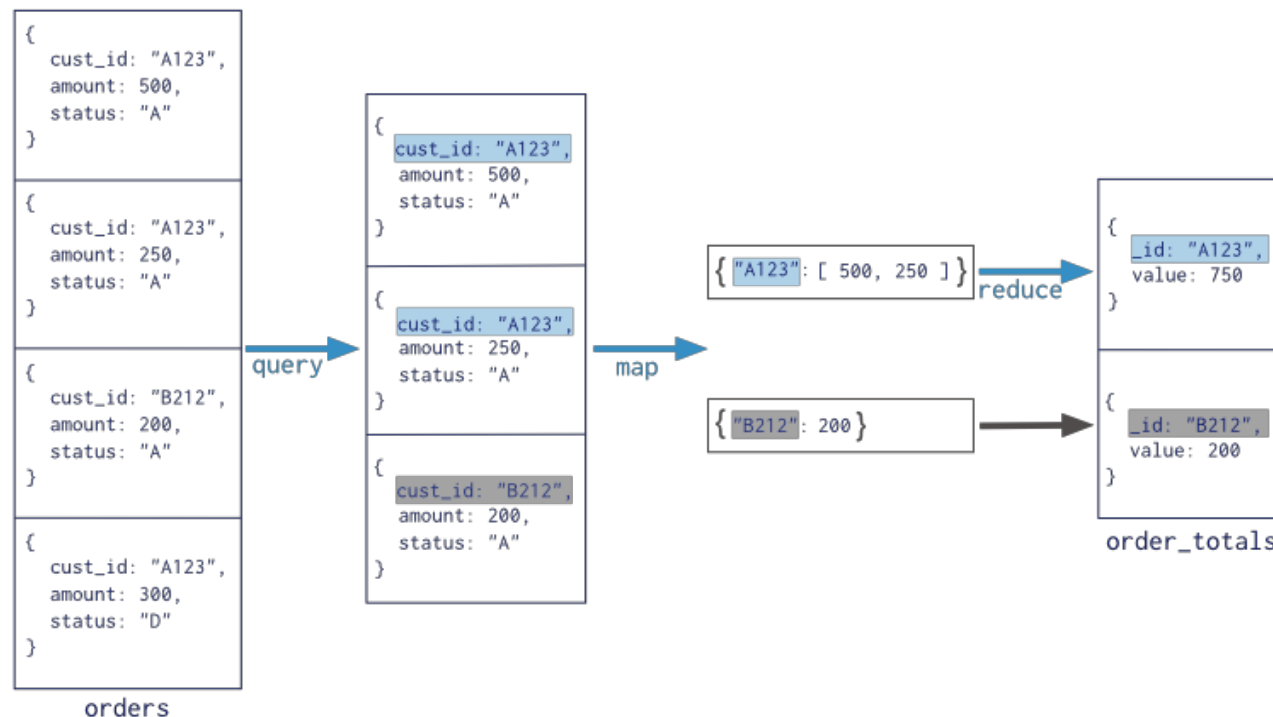
```
{ "cursor" : "BtreeCursor
display_1",
  "nscanned" : 1,
  "nscannedObjects" : 1,
  "n" : 1,
  "millis" : 0,
  "indexBounds" : {
    "display" : [ [ "+1 800-5650001",
"+1 800-5650001" ] ] }
}
```

# MapReduce processing

Collection  
↓  
db.orders.mapReduce(  
  map    → function() { emit( this.cust\_id, this.amount ); },  
  reduce → function(key, values) { return Array.sum( values ) },  
  query  → {  
  output →    query: { status: "A" },  
              out: "order\_totals"  
              }  
)

Execution order:

1. Query
2. Map
3. Reduce
4. [finalize] to wrap up reducer results, e.g. to take the max among all
5. [output]



# Replica sets

**Duplication (replication) to prevent against server failure and data loss**

Example (three servers): mkdir ./mongo1 ./mongo2 ./mongo3

Create replication set:

```
mongod --replSet book --dbpath ./mongo1 --port 27011 --rest
mongod --replSet book --dbpath ./mongo2 --port 27012 --rest
mongod --replSet book --dbpath ./mongo3 --port 27013 --rest
```

Then in one of the servers initialize replication set:

```
mongo localhost:27011
```

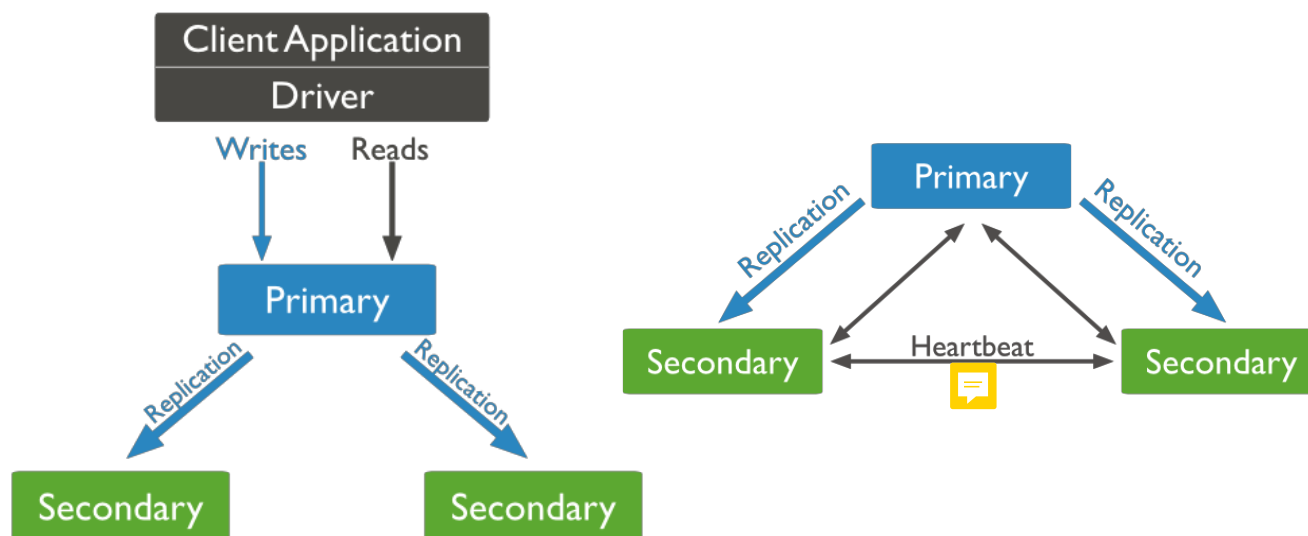
```
> rs.initiate({_id: 'book',
               members: [ {_id: 1, host: 'localhost:27011'},
                           {_id: 2, host: 'localhost:27012'},
                           {_id: 3, host: 'localhost:27013'} ] })
```

Then one server will output **[rs Manager] replSet PRIMARY**  
while two will output **[rs sync] replSet SECONDARY**

# Replica sets in MongoDB

The servers held a vote to determine who is the master (primary); the two others are replicas ("secondary")

By default, applications *read/write through the primary*, who pushes updates to secondary servers

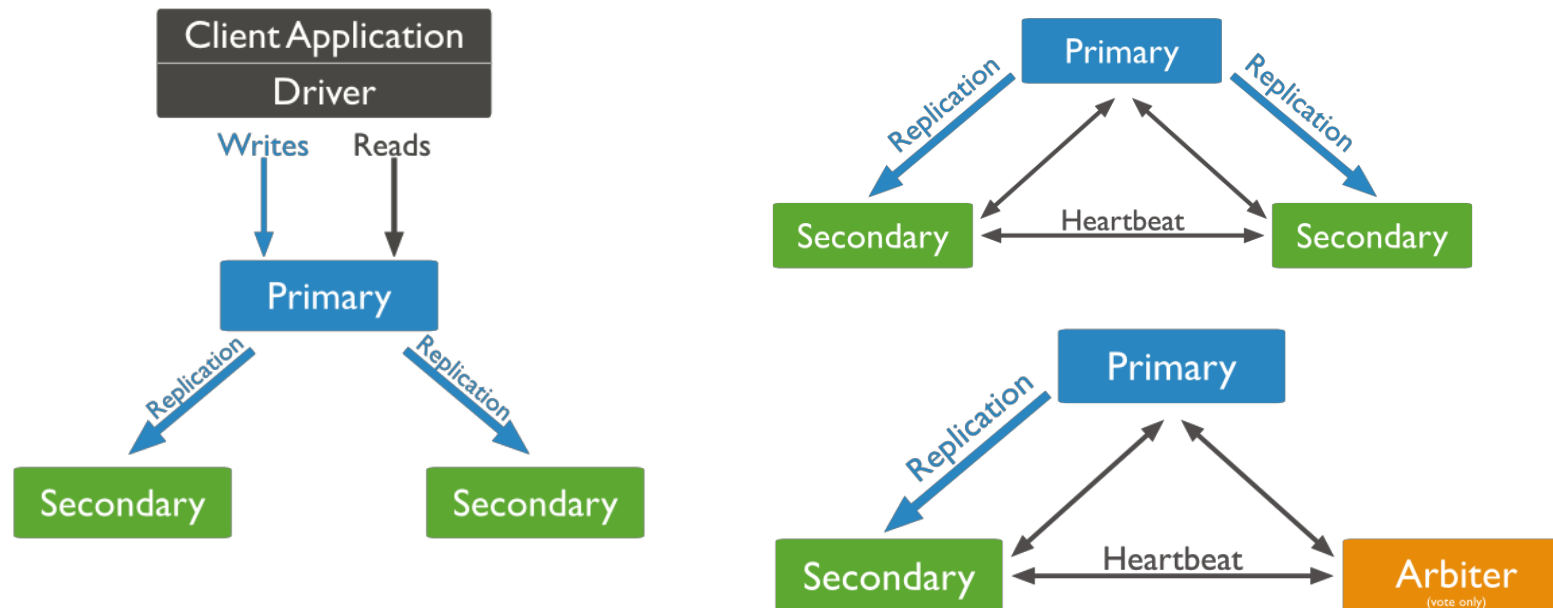


If the master is considered dead, there are new elections

- *Only succeed if more than half of the original replication set votes*
- Operations attempted on a demoted (dead) master are lost
- *A write is considered successful only if > half of the replicas "saw" it*

# Replica sets

By default, write to the primary, who pushes updates to secondary servers



MongoDB recommends an *odd number of servers in a replica set*, to allow a majority in case of network failure. **One arbiter** may be added to the replica set

- Strong consistency on read
- Resistance to some partitioning

# MongoDB sharding

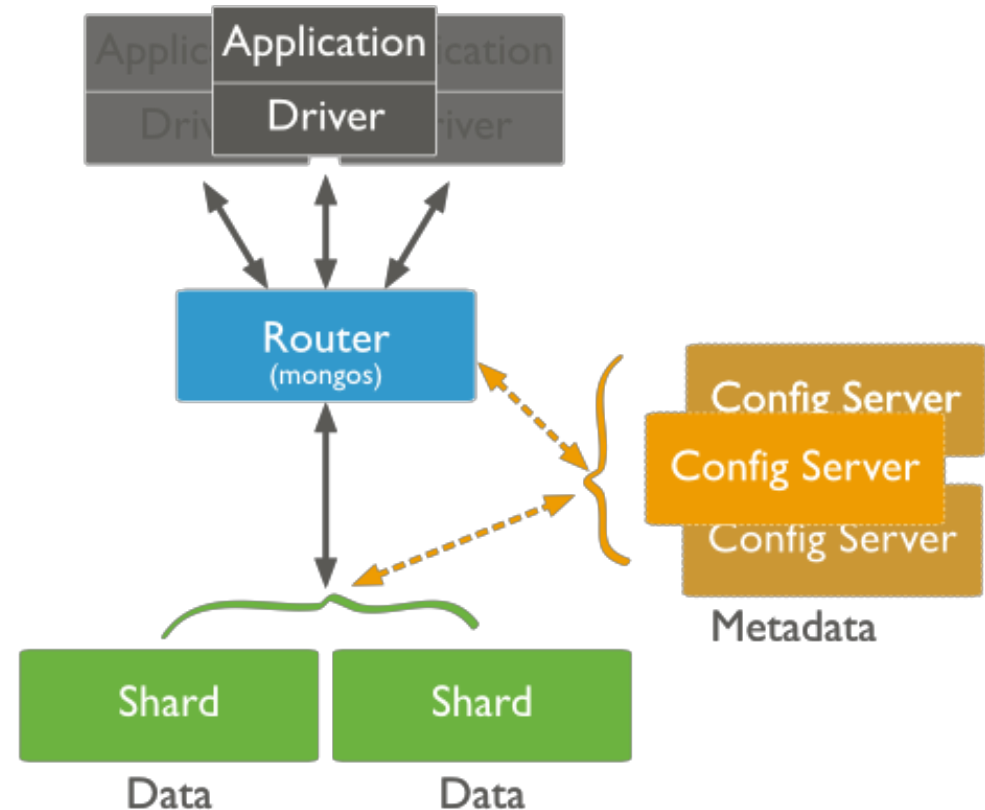
*Sharding* = partitioning

1 shard = 1 fragment

- To distribute a very large collection across several servers

Sharding is logically *on top of replication*

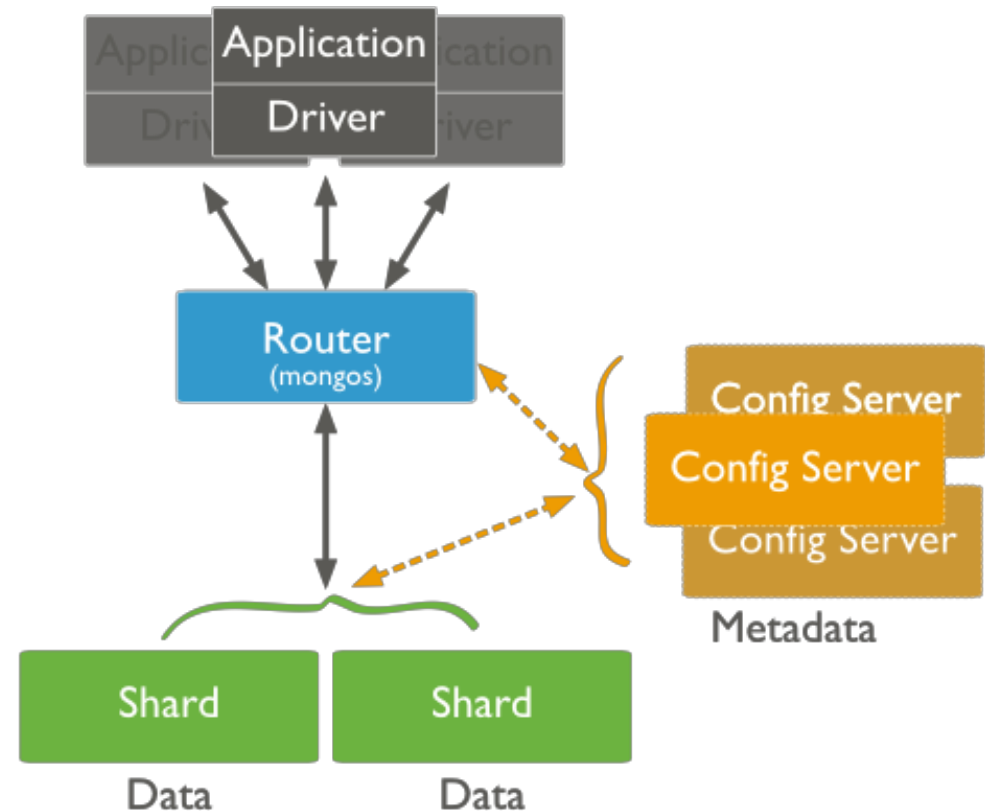
- Each shard server may participate to a replica set



# MongoDB sharding

## Roles:

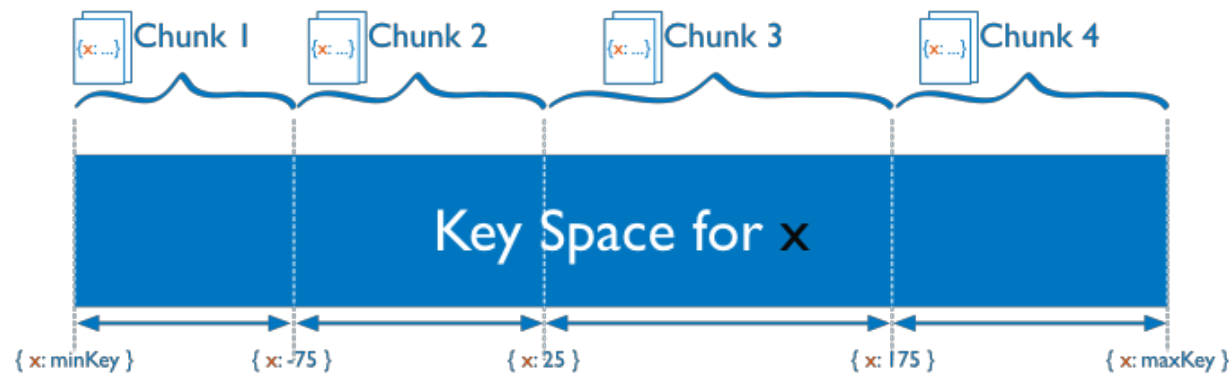
- **Shard** (shard server): stores a collection fragment
- **Config server(s)**: store(s) information on which shard has what (~ catalog)
- **Single point of entry**: mongos





# MongoDB sharding

- A data collection is partitioned into **chunks** based on the value of a **shard key**
- Each chunk covers a key range



- 1 shard = a set of chunks
- Mongos routes writes to the appropriate chunk based on the shard key value
- Chunks are split when they grow beyond a fixed chunk size (64 MB default, can be split)
- MongoDB migrates chunks across shards for load balancing

# MongoDB sharding

*// starting the shard servers:*

```
mkdir ./mongo4 ./mongo5
```

```
mongod --shardsvr --dbpath ./mongo4 --port 27014
```

```
mongod --shardsvr --dbpath ./mongo5 --port 27015
```

*// starting the config server:*

```
mkdir ./mongoconfig
```

```
mongod --configsvr --dbpath ./mongoconfig --port 27016
```

*// starting mongos connected to the config*

```
mongos --configdb localhost:27016 --chunkSize 1 --port 27020
```

*// talking to mongos to configure sharding:*

```
mongo localhost:27020/admin
```

```
> db.runCommand( { addshard : "localhost:27014" } ) -> { "shardAdded" :  
"shard0000", "ok" : 1 }
```

```
> db.runCommand( { addshard : "localhost:27015" } ) -> { "shardAdded" :  
"shard0001", "ok" : 1 }
```

```
> db.runCommand({ enablesharding: "test"}) -> { "ok" : 1 }
```

```
> db.runCommand( { shardcollection : "test.cities", key : {name : 1} } ) //  
{ "collectionsharded" : "test.cities", "ok" : 1 }
```

# Another JSON store: CouchDB



<http://db-engines.com/en/system/CouchDB%3BMongoDB>

| Feature                       | CouchDB                                               | MongoDB                                       |
|-------------------------------|-------------------------------------------------------|-----------------------------------------------|
| Since                         | 2005                                                  | 2009                                          |
| Ranking                       | #23 overall, #2 document store                        | #4 overall, #1 document store                 |
| From                          | Apache Software                                       | MongoDB Inc                                   |
| APIs and other access methods | RESTful HTTP/JSON API                                 | proprietary protocol using JSON               |
| Replication methods           | Master-master replication<br>Master-slave replication | Master-slave replication                      |
| MapReduce                     | yes                                                   | yes                                           |
| Consistency concepts          | Eventual Consistency                                  | Eventual Consistency<br>Immediate Consistency |
| Foreign keys                  | no                                                    | no                                            |

More document stores: [http://db-engines.com/en/ranking\\_trend/document+store](http://db-engines.com/en/ranking_trend/document+store)

# Concurrency control in CouchDB

- Update granularity = document
- To change a document's attribute, rewrite the document!
- Multi-Version Concurrency Control: some requests may return "old" versions but they each return a version that was valid at some point




# Synchronization in CouchDB

Each server can work independently

**Incremental replication** can be set to run in the background

**Synchronization** is on demand between any pair of servers

Diverging changes are flagged as **conflicts**; conflict resolution policy must be specified. One document version wins, the other is considered older.



# Conclusion: JSON stores

- **JSON**: extremely popular data interchange format
- **MongoDB**: (by far) most popular JSON data management system
- **MongoDB query language**:
  - Rich matching within one document
  - This is declarative.
  - Pipeline processing for more complex operations
  - Non-declarative
- **MongoDB replication**:
  - replica group, voting, quorum
- **MongoDB distribution**: sharding
- Other JSON stores:
  - Amazon DynamoDB (also serves as K-V store)
  - CouchDB